

# 操作系统原理

## 第二章：计算机系统结构

洪明坚

重庆大学软件学院

February 19, 2016

- 1 Computer system structure
  - Bootstrap
  - Interrupt and Exception
  - I/O structure
  - Hardware protection
  - System call
  - Function calling convention in C

# Computer system structure

# Computer system structure

- We need to have a general knowledge of the structure of a computer system before we can explore the details of the operating system.

# Computer system structure

- We need to have a general knowledge of the structure of a computer system before we can explore the details of the operating system.
  - You should be familiar with these concepts in the course *Computer architecture*.

# Computer system structure

- We need to have a general knowledge of the structure of a computer system before we can explore the details of the operating system.
  - You should be familiar with these concepts in the course *Computer architecture*.
  - But, some of them will be explored in-depth with the operating system in mind.

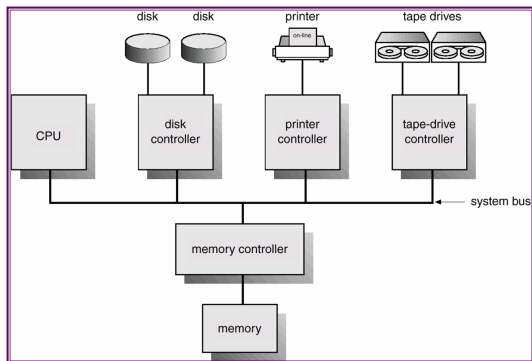
# Computer system structure

- We need to have a general knowledge of the structure of a computer system before we can explore the details of the operating system.
  - You should be familiar with these concepts in the course *Computer architecture*.
  - But, some of them will be explored in-depth with the operating system in mind.

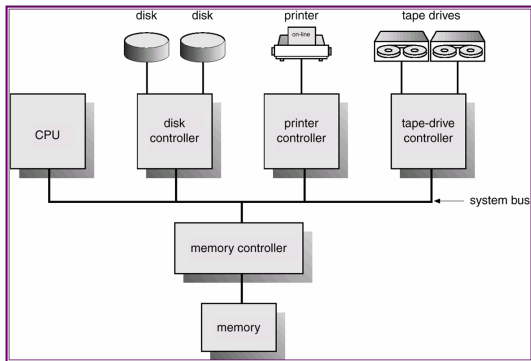
# Overall structure



# Overall structure

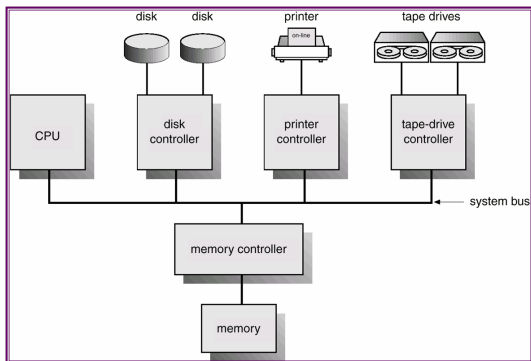


# Overall structure



- The CPU and device controllers can execute concurrently, competing for memory cycles.

# Overall structure



- The CPU and device controllers can execute concurrently, competing for memory cycles.
- To ensure orderly access to the memory, a memory controller is provided to synchronize access to the memory.



- When users just power on a computer, there is no already running operating system available.

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transfered to the entry of the operating system with a very basic environment.

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transfered to the entry of the operating system with a very basic environment.
  - This procedure is named **bootstrap** (or simply **boot**) an operating system.



# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transferred to the entry of the operating system with a very basic environment.
  - This procedure is named **bootstrap** (or simply **boot**) an operating system.
  - The program which boots the operating system is called **boot-loader**.

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transferred to the entry of the operating system with a very basic environment.
  - This procedure is named **bootstrap** (or simply **boot**) an operating system.
  - The program which boots the operating system is called **boot-loader**.
- Examples
  - **NTLDR** - boot-loader for Windows NT/XP (resides in C:\).

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transferred to the entry of the operating system with a very basic environment.
  - This procedure is named **bootstrap** (or simply **boot**) an operating system.
  - The program which boots the operating system is called **boot-loader**.
- Examples
  - **NTLDR** - boot-loader for Windows NT/XP (resides in C:\).
  - **GRUB** - one of boot-loaders for the Unix/Linux.

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transferred to the entry of the operating system with a very basic environment.
  - This procedure is named **bootstrap** (or simply **boot**) an operating system.
  - The program which boots the operating system is called **boot-loader**.
- Examples
  - **NTLDR** - boot-loader for Windows NT/XP (resides in C:\).
  - **GRUB** - one of boot-loaders for the Unix/Linux.
- Bear in mind that

# Bootstrap

- When users just power on a computer, there is no already running operating system available.
  - We must load the operating system kernel from some persistent storages, such as disk and network server, to memory.
  - Then the control is transferred to the entry of the operating system with a very basic environment.
  - This procedure is named **bootstrap** (or simply **boot**) an operating system.
  - The program which boots the operating system is called **boot-loader**.
- Examples
  - **NTLDR** - boot-loader for Windows NT/XP (resides in C:\).
  - **GRUB** - one of boot-loaders for the Unix/Linux.
- Bear in mind that
  - **The boot-loader is NOT part of the operating system.**

# Questions

- Any questions?



# Interrupt

# Interrupt

- Modern computers and operating systems are **interrupt driven**.



# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.

# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.
- When the CPU is interrupted, it must **serve the interrupt** by

# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.
- When the CPU is interrupted, it must **serve the interrupt** by
  - ① *Hardware*: saves some of registers and branches to the **interrupt service routine (ISR)**;

# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.
- When the CPU is interrupted, it must **serve the interrupt** by
  - ① *Hardware*: saves some of registers and branches to the **interrupt service routine (ISR)**;
  - ② *Assembly language procedure in ISR*: saves rest of registers if necessary and sets up a convenient environment;

# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.
- When the CPU is interrupted, it must **serve the interrupt** by
  - ① *Hardware*: saves some of registers and branches to the **interrupt service routine (ISR)**;
  - ② *Assembly language procedure in ISR*: saves rest of registers if necessary and sets up a convenient environment;
  - ③ *C language procedure in ISR*: does serve the interrupt, typically reads and buffers input data from peripheral device;

# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.
- When the CPU is interrupted, it must **serve the interrupt** by
  - ① *Hardware*: saves some of registers and branches to the **interrupt service routine (ISR)**;
  - ② *Assembly language procedure in ISR*: saves rest of registers if necessary and sets up a convenient environment;
  - ③ *C language procedure in ISR*: does serve the interrupt, typically reads and buffers input data from peripheral device;
  - ④ *C language procedure in ISR*: returns to the *assembly language procedure in ISR*;

# Interrupt

- Modern computers and operating systems are **interrupt driven**.
  - Peripheral devices use **interrupt** to signal the CPU that something has happened.
- When the CPU is interrupted, it must **serve the interrupt** by
  - ① *Hardware*: saves some of registers and branches to the **interrupt service routine (ISR)**;
  - ② *Assembly language procedure in ISR*: saves rest of registers if necessary and sets up a convenient environment;
  - ③ *C language procedure in ISR*: does serve the interrupt, typically reads and buffers input data from peripheral device;
  - ④ *C language procedure in ISR*: returns to the *assembly language procedure in ISR*;
  - ⑤ *Assembly language procedure in ISR*: restores saved registers and returns to the location being interrupted.

# Interrupt Vector



# Interrupt Vector

- Usually, a computer system has several peripheral devices.

# Interrupt Vector

- Usually, a computer system has several peripheral devices.
- When an interrupt occurs, CPU must know which device triggered it.

# Interrupt Vector

- Usually, a computer system has several peripheral devices.
- When an interrupt occurs, CPU must know which device triggered it.
  - Computer system assigns each device an *unique* interrupt request number (e.g., an 8-bit integer), or simply **IRQ**.

# Interrupt Vector

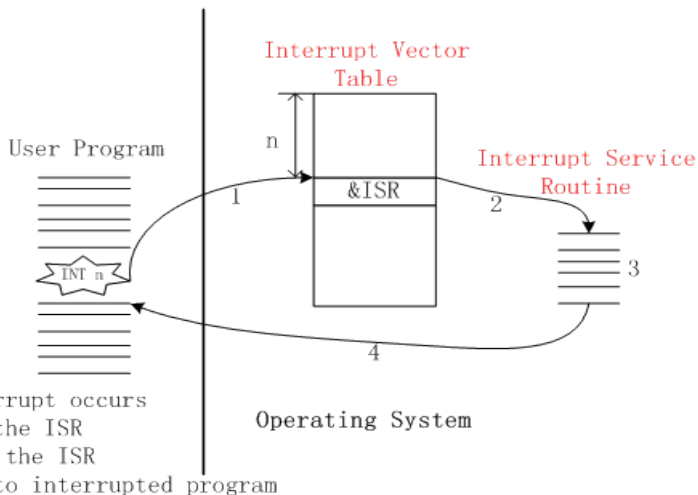
- Usually, a computer system has several peripheral devices.
- When an interrupt occurs, CPU must know which device triggered it.
  - Computer system assigns each device an *unique* interrupt request number (e.g., an 8-bit integer), or simply **IRQ**.
  - The addresses of all ISRs are collected into a table called **interrupt vector**.

# Interrupt Vector

- Usually, a computer system has several peripheral devices.
- When an interrupt occurs, CPU must know which device triggered it.
  - Computer system assigns each device an *unique* interrupt request number (e.g., an 8-bit integer), or simply **IRQ**.
  - The addresses of all ISRs are collected into a table called **interrupt vector**.
  - When servicing an interrupt, CPU uses IRQ to index the interrupt vector to fetch the address of ISR and branches to it.

# Put them all together

# Put them all together



1. An interrupt occurs
2. Locate the ISR
3. Execute the ISR
4. Return to interrupted program

# Exception



# Exception

- Interrupt

# Exception

- Interrupt
  - Triggered by peripheral devices;

# Exception

- Interrupt
  - Triggered by peripheral devices;
  - Asynchronous.

# Exception

- Interrupt
  - Triggered by peripheral devices;
  - Asynchronous.
- **Exception**

# Exception

- Interrupt
  - Triggered by peripheral devices;
  - Asynchronous.
- **Exception**
  - Exceptions occur when the processor detects an error condition while executing an instruction, such as division by zero and invalid memory access.

# Exception

- Interrupt

- Triggered by peripheral devices;
- Asynchronous.

- **Exception**

- Exceptions occur when the processor detects an error condition while executing an instruction, such as division by zero and invalid memory access.
- Synchronous.

# Exception

- Interrupt
  - Triggered by peripheral devices;
  - Asynchronous.
- **Exception**
  - Exceptions occur when the processor detects an error condition while executing an instruction, such as division by zero and invalid memory access.
  - Synchronous.
- Other than the above, handling of interrupts and exceptions is identical.

# Exception

- Interrupt
  - Triggered by peripheral devices;
  - Asynchronous.
- **Exception**
  - Exceptions occur when the processor detects an error condition while executing an instruction, such as division by zero and invalid memory access.
  - Synchronous.
- Other than the above, handling of interrupts and exceptions is identical.
  - Exception is also known as **software-generated interrupt** or **synchronous interrupt**.



# Questions

- Any questions?



# I/O structure

---

<sup>1</sup>Also known as blocking or non-overlapping I/O.

<sup>2</sup>Also known as non-blocking or overlapping I/O.

# I/O structure

- When CPU is doing I/O with a peripheral device, two methods is available:

---

<sup>1</sup>Also known as blocking or non-overlapping I/O.

<sup>2</sup>Also known as non-blocking or overlapping I/O.

# I/O structure

- When CPU is doing I/O with a peripheral device, two methods is available:
  - (a) synchronous <sup>1</sup>

---

<sup>1</sup>Also known as blocking or non-overlapping I/O.

<sup>2</sup>Also known as non-blocking or overlapping I/O.

# I/O structure

- When CPU is doing I/O with a peripheral device, two methods is available:
  - (a) synchronous <sup>1</sup> and (b) asynchronous <sup>2</sup> I/O.

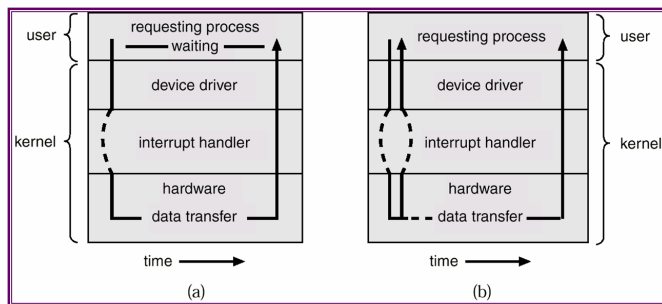
---

<sup>1</sup>Also known as blocking or non-overlapping I/O.

<sup>2</sup>Also known as non-blocking or overlapping I/O.

# I/O structure

- When CPU is doing I/O with a peripheral device, two methods is available:
  - (a) synchronous<sup>1</sup> and (b) asynchronous<sup>2</sup> I/O.



<sup>1</sup>Also known as blocking or non-overlapping I/O.

<sup>2</sup>Also known as non-blocking or overlapping I/O.

# How CPU accesses peripheral devices?

# How CPU accesses peripheral devices?

- CPU accesses devices through device controllers.



# How CPU accesses peripheral devices?

- CPU accesses devices through device controllers.
  - Device controllers include registers to hold commands and the data being transferred.

# How CPU accesses peripheral devices?

- CPU accesses devices through device controllers.
  - Device controllers include registers to hold commands and the data being transferred.
  - How CPU accesses these registers?

# How CPU accesses peripheral devices?

- CPU accesses devices through device controllers.
  - Device controllers include registers to hold commands and the data being transferred.
  - How CPU accesses these registers?
- Two methods:

# How CPU accesses peripheral devices?

- CPU accesses devices through device controllers.
  - Device controllers include registers to hold commands and the data being transferred.
  - How CPU accesses these registers?
- Two methods:
  - I/O port

# How CPU accesses peripheral devices?

- CPU accesses devices through device controllers.
  - Device controllers include registers to hold commands and the data being transferred.
  - How CPU accesses these registers?
- Two methods:
  - I/O port
  - Memory-mapped I/O

# I/O port (1/2)

# I/O port (1/2)

- All registers within device controllers are collected.

# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.



# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.
- Special I/O instructions are designed to allow data transfers between these registers and memory.

# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.
- Special I/O instructions are designed to allow data transfers between these registers and memory.
- Example: IBM-PC

# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.
- Special I/O instructions are designed to allow data transfers between these registers and memory.
- Example: IBM-PC
  - 16-bit I/O ports are used to address the registers of device controllers.

# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.
- Special I/O instructions are designed to allow data transfers between these registers and memory.
- Example: IBM-PC
  - 16-bit I/O ports are used to address the registers of device controllers.
  - Two special I/O instructions: **IN** and **OUT** are included in the INTEL x86 CPU.

# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.
- Special I/O instructions are designed to allow data transfers between these registers and memory.
- Example: IBM-PC
  - 16-bit I/O ports are used to address the registers of device controllers.
  - Two special I/O instructions: **IN** and **OUT** are included in the INTEL x86 CPU.
    - *IN reg, port* - Read a byte/word from *port* to CPU register *reg*.

# I/O port (1/2)

- All registers within device controllers are collected.
  - An unique address (which is called **port**, an 8- or 16-bit integer) is assigned to each of them.
- Special I/O instructions are designed to allow data transfers between these registers and memory.
- Example: IBM-PC
  - 16-bit I/O ports are used to address the registers of device controllers.
  - Two special I/O instructions: **IN** and **OUT** are included in the INTEL x86 CPU.
    - IN *reg, port* - Read a byte/word from *port* to CPU register *reg*.
    - OUT *port, reg* - Write the content of CPU register *reg* to *port*.

# I/O port (2/2)

## I/O port (2/2)

- Part of PC I/O port address map



# I/O port (2/2)

- Part of PC I/O port address map

Range (hex)	Function
000-01F	1st DMA controller
020-03F	1st Programmable Interrupt Controller (PIC)
040-05F	Programmable Interval Timer (System timer)
060-06F	Keyboard
220-233	Sound card
3D0-3DF	Color Graphics Adapter

# I/O port (2/2)

- Part of PC I/O port address map

Range (hex)	Function
000-01F	1st DMA controller
020-03F	1st Programmable Interrupt Controller (PIC)
040-05F	Programmable Interval Timer (System timer)
060-06F	Keyboard
220-233	Sound card
3D0-3DF	Color Graphics Adapter

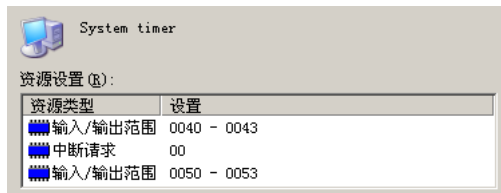
- Example: system timer

## I/O port (2/2)

- Part of PC I/O port address map

Range (hex)	Function
000-01F	1st DMA controller
020-03F	1st Programmable Interrupt Controller (PIC)
040-05F	Programmable Interval Timer (System timer)
060-06F	Keyboard
220-233	Sound card
3D0-3DF	Color Graphics Adapter

- Example: system timer



# Memory-mapped I/O (1/3)

# Memory-mapped I/O (1/3)

- In the method of I/O port,

# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.

# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.
- Registers within device controller is just a piece of storage.

# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.
- Registers within device controller is just a piece of storage.
  - Why not access these registers using the same method as memory?



# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.
- Registers within device controller is just a piece of storage.
  - Why not access these registers using the same method as memory?
  - In this case, a (unique) **memory address** is assigned to every register, **NOT a port address**.

# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.
- Registers within device controller is just a piece of storage.
  - Why not access these registers using the same method as memory?
  - In this case, a (unique) **memory address** is assigned to every register, **NOT a port address**.
- That's the memory-mapped I/O.

# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.
- Registers within device controller is just a piece of storage.
  - Why not access these registers using the same method as memory?
  - In this case, a (unique) **memory address** is assigned to every register, **NOT a port address**.
- That's the memory-mapped I/O.
  - Memory-mapped I/O uses the same bus to address both memory and I/O devices

# Memory-mapped I/O (1/3)

- In the method of I/O port,
  - we can view I/O ports as another separate address space, independent of memory address space.
- Registers within device controller is just a piece of storage.
  - Why not access these registers using the same method as memory?
  - In this case, a (unique) **memory address** is assigned to every register, **NOT a port address**.
- That's the memory-mapped I/O.
  - Memory-mapped I/O uses the same bus to address both memory and I/O devices
  - In order to accommodate the I/O devices, areas of CPU addressable space must be reserved for I/O rather than memory.

# Memory-mapped I/O (2/3)

# Memory-mapped I/O (2/3)

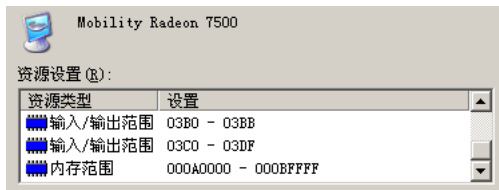
- Example: video controller of IBM-PC

# Memory-mapped I/O (2/3)

- Example: video controller of IBM-PC
  - Each location on the screen is mapped to a memory location.

# Memory-mapped I/O (2/3)

- Example: video controller of IBM-PC
  - Each location on the screen is mapped to a memory location.





# Memory-mapped I/O (3/3)

# Memory-mapped I/O (3/3)

- Advantages

# Memory-mapped I/O (3/3)

- Advantages
  - Every instruction that can reference memory can also reference device controller registers.

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

- Disadvantages

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

- Disadvantages

- Most computers nowadays have some form of caching memory words.  
But,

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

- Disadvantages

- Most computers nowadays have some form of caching memory words. But, caching a device controller register would be disastrous.



# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

- Disadvantages

- Most computers nowadays have some form of caching memory words. But, caching a device controller register would be disastrous.
- Modern computer systems use both of them,

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

- Disadvantages

- Most computers nowadays have some form of caching memory words. But, caching a device controller register would be disastrous.

- Modern computer systems use both of them,

- with memory-mapped I/O for data buffers and separate I/O ports for the command registers,

# Memory-mapped I/O (3/3)

- Advantages

- Every instruction that can reference memory can also reference device controller registers.
  - The device drivers can be written entirely in C.
- No special protection mechanism is needed to keep user processes from performing I/O.

- Disadvantages

- Most computers nowadays have some form of caching memory words. But, caching a device controller register would be disastrous.

- Modern computer systems use both of them,

- with memory-mapped I/O for data buffers and separate I/O ports for the command registers,
- as in the previous example of *Mobility Radeon 7500*.

# Questions

- Any questions?



# Questions

- Any questions?



# Hardware protection

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways



# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation**

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.
  - **Privileged instructions**

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.
  - **Privileged instructions** - Prevent user programs disrupting the normal operation of the system by issuing illegal I/O instructions.

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.
  - **Privileged instructions** - Prevent user programs disrupting the normal operation of the system by issuing illegal I/O instructions.
  - **Memory protection**

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.
  - **Privileged instructions** - Prevent user programs disrupting the normal operation of the system by issuing illegal I/O instructions.
  - **Memory protection** - Prevent a user program directly accessing the memory of another user program or even operating system.

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.
  - **Privileged instructions** - Prevent user programs disrupting the normal operation of the system by issuing illegal I/O instructions.
  - **Memory protection** - Prevent a user program directly accessing the memory of another user program or even operating system.
  - **CPU protection**

# Hardware protection

- To ensure proper operation, we must protect the operating system and all other program and their data from any malfunctioning program.
- Hardware protection can be broken down different ways
  - **Dual-mode operation** - Prevent user programs taking over part of the OS and using this to overwrite other programs or even modify the OS itself.
  - **Privileged instructions** - Prevent user programs disrupting the normal operation of the system by issuing illegal I/O instructions.
  - **Memory protection** - Prevent a user program directly accessing the memory of another user program or even operating system.
  - **CPU protection** - Prevent a user program from getting stuck in an infinite loop and never returning control to the operating system.



# Dual-mode operation

# Dual-mode operation

- We need at least two separate modes of operation:

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode**

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode**

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.
- **Mode bit** is added to computer hardware to indicate the current mode: monitor (0) or user (1).



# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.
- **Mode bit** is added to computer hardware to indicate the current mode: monitor (0) or user (1).
  - It's set to *monitor* at system boot time.

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.
- **Mode bit** is added to computer hardware to indicate the current mode: monitor (0) or user (1).
  - It's set to *monitor* at system boot time. The operating system is then loaded, and starts user programs in user mode.

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.
- **Mode bit** is added to computer hardware to indicate the current mode: monitor (0) or user (1).
  - It's set to *monitor* at system boot time. The operating system is then loaded, and starts user programs in user mode.
- When an interrupt or exception occurs hardware switches to monitor mode.

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.
- **Mode bit** is added to computer hardware to indicate the current mode: monitor (0) or user (1).
  - It's set to *monitor* at system boot time. The operating system is then loaded, and starts user programs in user mode.
- When an interrupt or exception occurs hardware switches to monitor mode.
  - Whenever the operating system gains control of the computer, it's in monitor mode.

# Dual-mode operation

- We need at least two separate modes of operation:
  - **User mode** - Execution on behalf of user programs;
  - **Monitor mode** - Execution on behalf of operating system.
    - Also known as **supervisor, system, privileged or kernel mode**.
- **Mode bit** is added to computer hardware to indicate the current mode: monitor (0) or user (1).
  - It's set to *monitor* at system boot time. The operating system is then loaded, and starts user programs in user mode.
- When an interrupt or exception occurs hardware switches to monitor mode.
  - Whenever the operating system gains control of the computer, it's in monitor mode.
  - And the system always switches to user mode before passing control to a user program.

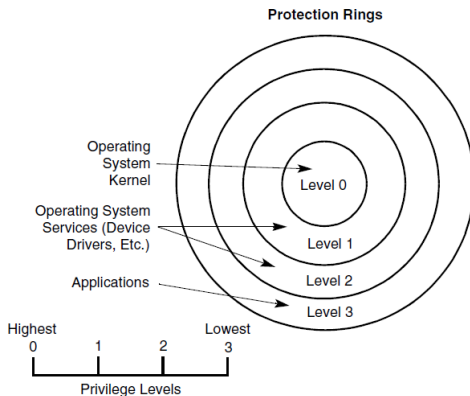
# Example

# Example

- INTEL IA-32 supports 4 modes to operate, named **protection rings**.

# Example

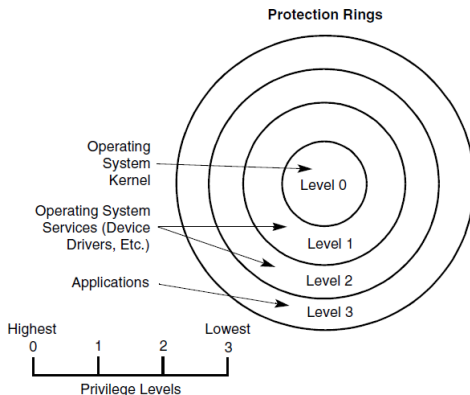
- INTEL IA-32 supports 4 modes to operate, named **protection rings**.





# Example

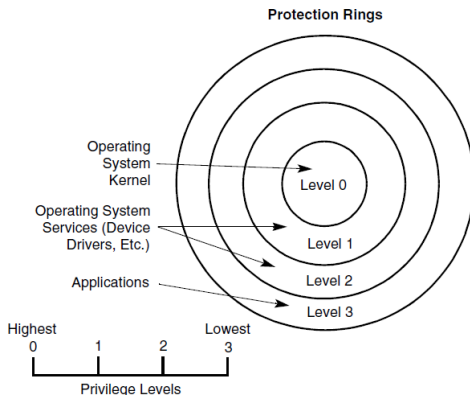
- INTEL IA-32 supports 4 modes to operate, named **protection rings**.



- But, most operating systems running on IA-32 only use 2 of 4.

# Example

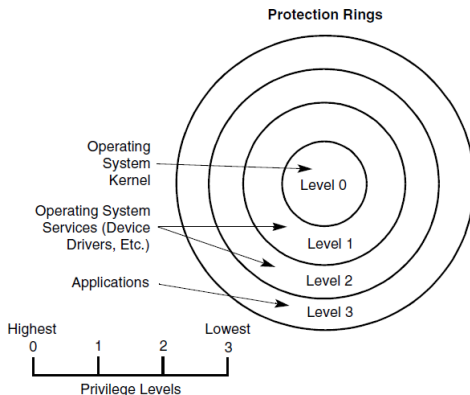
- INTEL IA-32 supports 4 modes to operate, named **protection rings**.



- But, most operating systems running on IA-32 only use 2 of 4.
  - Ring 0 as monitor mode;

## Example

- INTEL IA-32 supports 4 modes to operate, named **protection rings**.



- But, most operating systems running on IA-32 only use 2 of 4.
  - Ring 0 as monitor mode; ring 3 as user mode.

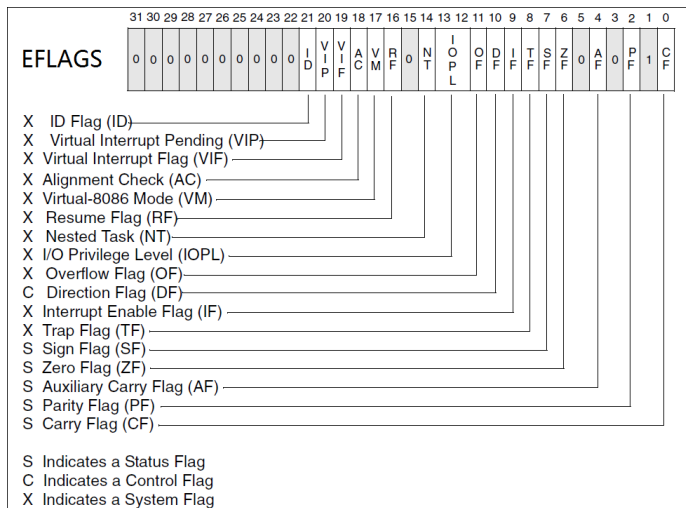
# Example (cont'd)

# Example (cont'd)

- **Mode bit** of INTEL IA-32

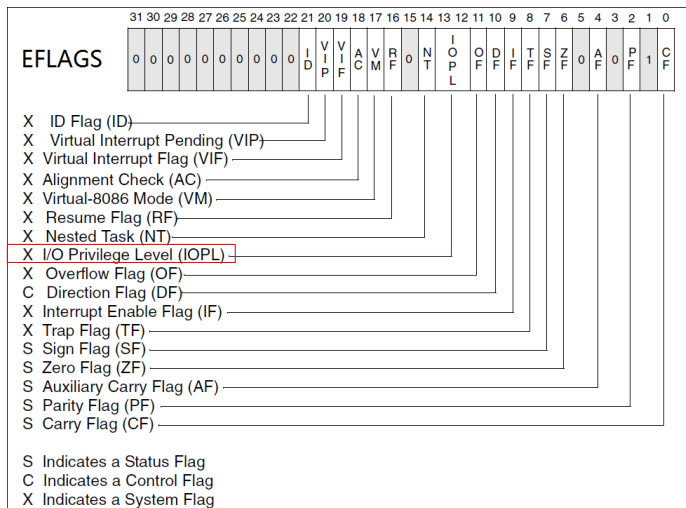
# Example (cont'd)

## • Mode bit of INTEL IA-32



# Example (cont'd)

## • Mode bit of INTEL IA-32



# I/O protection



# I/O protection

- All I/O instructions are privileged instructions.

# I/O protection

- All I/O instructions are privileged instructions.
  - The hardware allows privileged instructions to be executed only in monitor mode.

- All I/O instructions are privileged instructions.
  - The hardware allows privileged instructions to be executed only in monitor mode.
  - If these instructions are to be executed in user mode, the hardware does not execute the instruction, but rather treats it as illegal and generates an exception.

- All I/O instructions are privileged instructions.
  - The hardware allows privileged instructions to be executed only in monitor mode.
  - If these instructions are to be executed in user mode, the hardware does not execute the instruction, but rather treats it as illegal and generates an exception.
  - For example, **IN** and **OUT** are 2 privileged instructions in INTEL IA-32.

- All I/O instructions are privileged instructions.
  - The hardware allows privileged instructions to be executed only in monitor mode.
  - If these instructions are to be executed in user mode, the hardware does not execute the instruction, but rather treats it as illegal and generates an exception.
  - For example, **IN** and **OUT** are 2 privileged instructions in INTEL IA-32.
- Must ensure that a user program could never gain control of the computer in monitor mode.

# Memory protection

# Memory protection

- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:

# Memory protection

- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - **Base register**



# Memory protection

- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - **Base register** - holds the smallest legal physical memory address;

# Memory protection

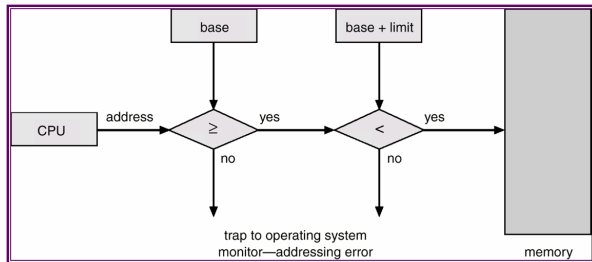
- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - **Base register** - holds the smallest legal physical memory address;
  - **Limit register**

# Memory protection

- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - **Base register** - holds the smallest legal physical memory address;
  - **Limit register** - contains the size of the range.

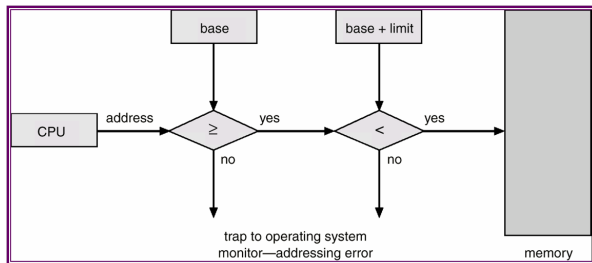
# Memory protection

- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - **Base register** - holds the smallest legal physical memory address;
  - **Limit register** - contains the size of the range.



# Memory protection

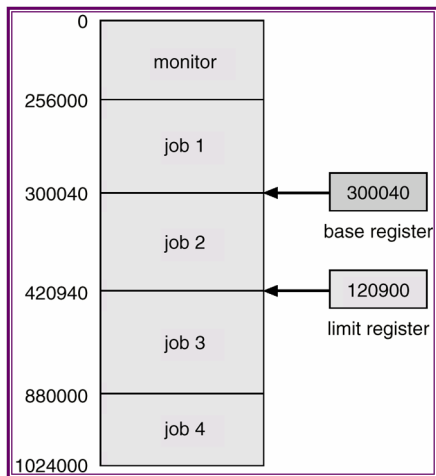
- In order to have memory protection, add two registers that determine the range of legal addresses a program may access:
  - **Base register** - holds the smallest legal physical memory address;
  - **Limit register** - contains the size of the range.



- We will return to this topic when entering *memory management*.

# Example

# Example



# CPU protection



- The operating system can enforce policies only if it gets a chance to run.

- The operating system can enforce policies only if it gets a chance to run.
  - If a malfunctioning program entered an infinite loop and never returns control to the operating system, then CPU was out of the control from the operating system.

- The operating system can enforce policies only if it gets a chance to run.
  - If a malfunctioning program entered an infinite loop and never returns control to the operating system, then CPU was out of the control from the operating system.
- **Timer**

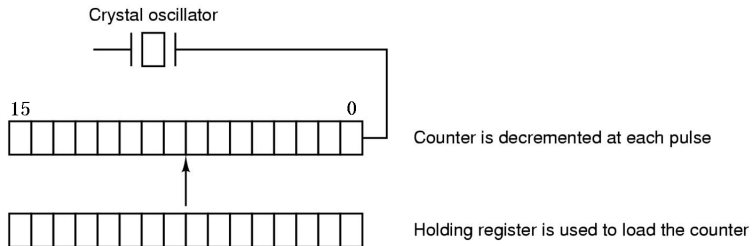
- The operating system can enforce policies only if it gets a chance to run.
  - If a malfunctioning program entered an infinite loop and never returns control to the operating system, then CPU was out of the control from the operating system.
- **Timer** - interrupts CPU after a specified period to ensure operating system maintains control.

- The operating system can enforce policies only if it gets a chance to run.
  - If a malfunctioning program entered an infinite loop and never returns control to the operating system, then CPU was out of the control from the operating system.
- **Timer** - interrupts CPU after a specified period to ensure operating system maintains control.
  - Remember that when interrupt occurs, the operating system will get the control via ISR.

- The operating system can enforce policies only if it gets a chance to run.
  - If a malfunctioning program entered an infinite loop and never returns control to the operating system, then CPU was out of the control from the operating system.
- **Timer** - interrupts CPU after a specified period to ensure operating system maintains control.
  - Remember that when interrupt occurs, the operating system will get the control via ISR.

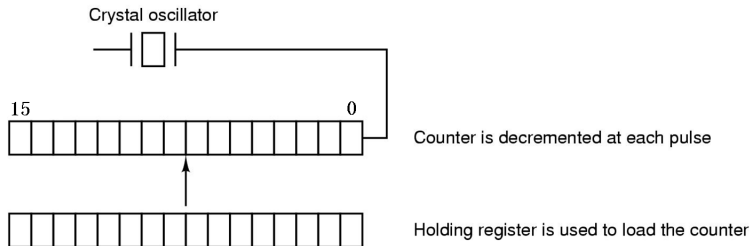
# Timer

# Timer



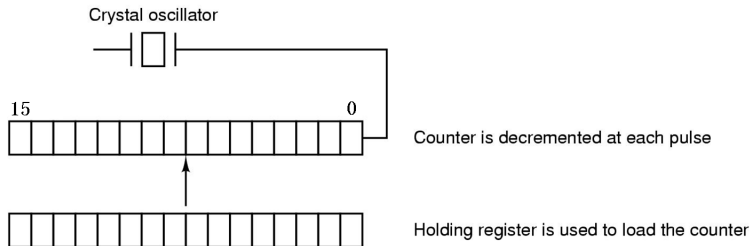


# Timer



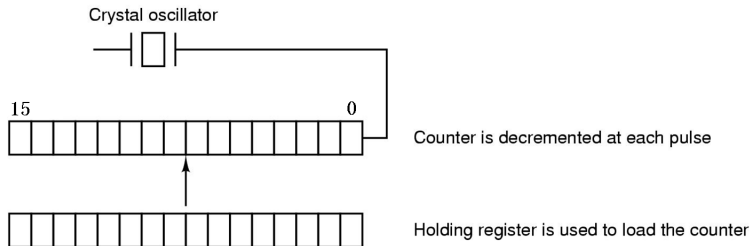
- The *counter register* is decremented by 1 at each pulse.

# Timer



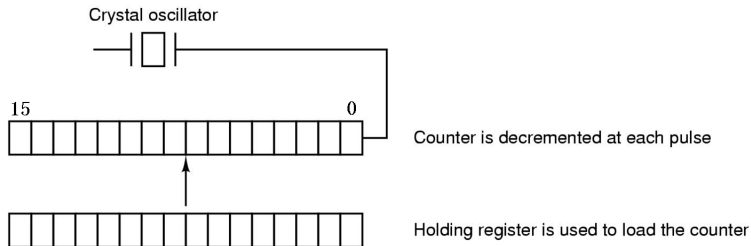
- The *counter register* is decremented by 1 at each pulse.
  - When the *counter register* reaches zero, the timer will interrupt CPU.

# Timer



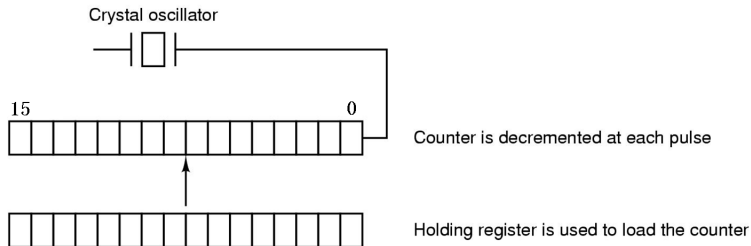
- The *counter register* is decremented by 1 at each pulse.
  - When the *counter register* reaches zero, the timer will interrupt CPU.
  - Then the *counter register* will be reloaded with the value of *holding register* and decrementing repeats.

# Timer



- The *counter register* is decremented by 1 at each pulse.
  - When the *counter register* reaches zero, the timer will interrupt CPU.
  - Then the *counter register* will be reloaded with the value of *holding register* and decrementing repeats.
- Example: Timer in IBM-PC

# Timer



- The *counter register* is decremented by 1 at each pulse.
  - When the *counter register* reaches zero, the timer will interrupt CPU.
  - Then the *counter register* will be reloaded with the value of *holding register* and decrementing repeats.
- Example: Timer in IBM-PC
  - An INTEL i8253 programmable interval timer with 16-bit counter and holding registers and pulses reach at 1193182Hz.

# Questions

- Any questions?



# System call

# System call

- The operating system does nothing useful itself.



# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.

# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.
  - How does operating system provide these services?

# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.
  - How does operating system provide these services?
- It's the **system call**

# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.
  - How does operating system provide these services?
- It's the **system call** - the (well-defined) **interface** between the operating system and the user programs.

# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.
  - How does operating system provide these services?
- It's the **system call** - the (well-defined) **interface** between the operating system and the user programs.
  - User programs can **ONLY** request services provided by the operating system via system call.

# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.
  - How does operating system provide these services?
- It's the **system call** - the (well-defined) **interface** between the operating system and the user programs.
  - User programs can **ONLY** request services provided by the operating system via system call.
  - The system calls in the interface vary from operating system to operating system.

# System call

- The operating system does nothing useful itself.
  - But it provides some useful services to the user programs, such as reading a file from disk and sending data to remote host via network adapter.
  - How does operating system provide these services?
- It's the **system call** - the (well-defined) **interface** between the operating system and the user programs.
  - User programs can **ONLY** request services provided by the operating system via system call.
  - The system calls in the interface vary from operating system to operating system.
  - Also known as **supervisor call**.

# Example



# Example

- Most operating systems provide the service that can read some bytes from a file: *read*

# Example

- Most operating systems provide the service that can read some bytes from a file: *read*
  - *count = read(fd, buffer, nbytes);*

# Example

- Most operating systems provide the service that can read some bytes from a file: *read*
  - $count = read(fd, buffer, nbytes);$
  - This system call reads *nbytes* data from file *fd* to the specified *buffer* and returns the number of bytes actually read in *count*.

# Procedure (1/2)

# Procedure (1/2)

1-3 Prepare parameters;

# Procedure (1/2)

- 1-3 Prepare parameters;
- 4 Call the wrapper (written in assembly language) of system call;

# Procedure (1/2)

- 1-3 Prepare parameters;
- 4 Call the wrapper (written in assembly language) of system call;
- 5 Store the **system call number** of *read* into a register;

# Procedure (1/2)

- 1-3 Prepare parameters;
- 4 Call the wrapper (written in assembly language) of system call;
- 5 Store the **system call number** of *read* into a register;
- 6 Trap into the operating system;



# Procedure (1/2)

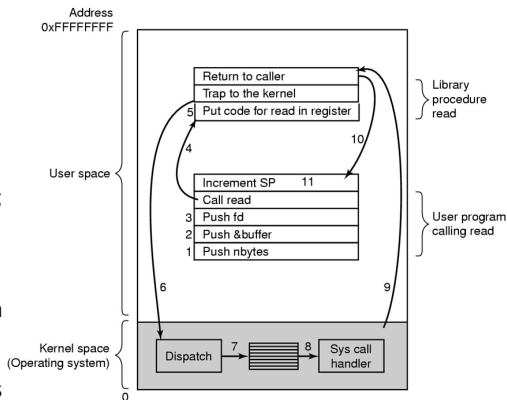
- 1-3 Prepare parameters;
- 4 Call the wrapper (written in assembly language) of system call;
- 5 Store the **system call number** of *read* into a register;
- 6 Trap into the operating system;
- 7 Get **system call service routine** for *read* by indexing a **system call table** using system call number;

# Procedure (1/2)

- 1-3 Prepare parameters;
- 4 Call the wrapper (written in assembly language) of system call;
- 5 Store the **system call number** of *read* into a register;
- 6 Trap into the operating system;
- 7 Get **system call service routine** for *read* by indexing a **system call table** using system call number;
- 8-11 System call service routine runs and returns to user programs on completion.

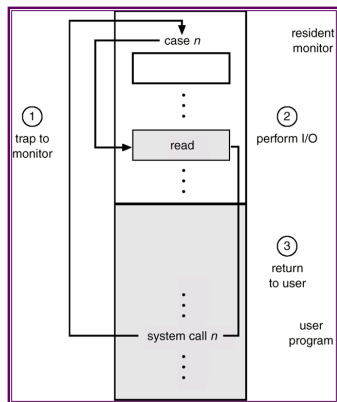
# Procedure (1/2)

- 1-3 Prepare parameters;
- 4 Call the wrapper (written in assembly language) of system call;
- 5 Store the **system call number** of *read* into a register;
- 6 Trap into the operating system;
- 7 Get **system call service routine** for *read* by indexing a **system call table** using system call number;
- 8-11 System call service routine runs and returns to user programs on completion.

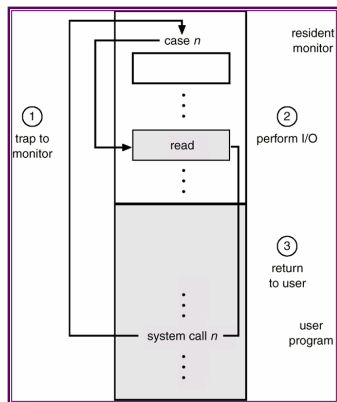


# Procedure (2/2)

# Procedure (2/2)

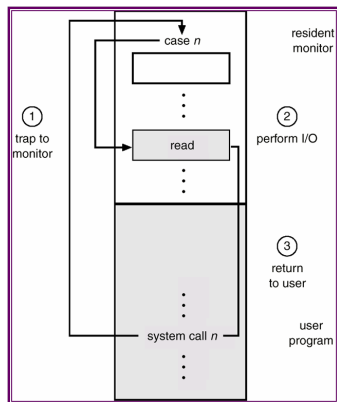


## Procedure (2/2)



- Resident monitor (or simply monitor) here means the operating system and

## Procedure (2/2)



- Resident monitor (or simply monitor) here means the operating system and
- $n$  is the system call number.

# Trap into the operating system (1/2)



# Trap into the operating system (1/2)

- The user programs cannot trap into the operating system directly.

# Trap into the operating system (1/2)

- The user programs cannot trap into the operating system directly.
- How to trap into the operating system?

# Trap into the operating system (1/2)

- The user programs cannot trap into the operating system directly.
- How to trap into the operating system?
  - Method 1: **Exception (software-generated interrupt)**.

# Trap into the operating system (1/2)

- The user programs cannot trap into the operating system directly.
- How to trap into the operating system?
  - Method 1: **Exception (software-generated interrupt)**.
  - Method 2: Special instruction.

# Trap into the operating system (2/2)

---

<sup>3</sup>Short for SoftWare Interrupt.

# Trap into the operating system (2/2)

- Exception

---

<sup>3</sup>Short for SoftWare Interrupt.

# Trap into the operating system (2/2)

- Exception
  - INTEL IA-32 provides an instruction to trigger a exception, *INT*.

---

<sup>3</sup>Short for SoftWare Interrupt.

# Trap into the operating system (2/2)

- Exception
  - INTEL IA-32 provides an instruction to trigger a exception, *INT*.
    - For example, Linux/FreeBSD uses *INT 0x80* to trap into the operating system and Windows NT/XP uses *INT 0x2e*.

---

<sup>3</sup>Short for SoftWare Interrupt.



# Trap into the operating system (2/2)

- Exception
  - INTEL IA-32 provides an instruction to trigger a exception, *INT*.
    - For example, Linux/FreeBSD uses *INT 0x80* to trap into the operating system and Windows NT/XP uses *INT 0x2e*.
- Special instruction

---

<sup>3</sup>Short for SoftWare Interrupt.

# Trap into the operating system (2/2)

- Exception
  - INTEL IA-32 provides an instruction to trigger a exception, *INT*.
    - For example, Linux/FreeBSD uses *INT 0x80* to trap into the operating system and Windows NT/XP uses *INT 0x2e*.
- Special instruction
  - In addition, INTEL IA-32 provides two special instructions to trap into the operating system: *SYSENTER* and *SYSEXIT* because of the extra overhead of *INT* instruction.

---

<sup>3</sup>Short for SoftWare Interrupt.

# Trap into the operating system (2/2)

- Exception

- INTEL IA-32 provides an instruction to trigger a exception, *INT*.
  - For example, Linux/FreeBSD uses *INT 0x80* to trap into the operating system and Windows NT/XP uses *INT 0x2e*.

- Special instruction

- In addition, INTEL IA-32 provides two special instructions to trap into the operating system: *SYSENTER* and *SYSEXIT* because of the extra overhead of *INT* instruction.
  - Only supported on processors after Pentium II, i.e., Family 6, Model 3, Stepping 3.

---

<sup>3</sup>Short for SoftWare Interrupt.

# Trap into the operating system (2/2)

- Exception

- INTEL IA-32 provides an instruction to trigger a exception, *INT*.
  - For example, Linux/FreeBSD uses *INT 0x80* to trap into the operating system and Windows NT/XP uses *INT 0x2e*.

- Special instruction

- In addition, INTEL IA-32 provides two special instructions to trap into the operating system: *SYSENTER* and *SYSEXIT* because of the extra overhead of *INT* instruction.
  - Only supported on processors after Pentium II, i.e., Family 6, Model 3, Stepping 3.
- ARM processors use *swi*<sup>3</sup> to trap into the operating system.

---

<sup>3</sup>Short for SoftWare Interrupt.

# System call v.s. library function

# System call v.s. library function

- System call will trap into the OS kernel; while library function does not.

# System call v.s. library function

- System call will trap into the OS kernel; while library function does not.
  - So, system call is MUCH more slow than library function.

# System call v.s. library function

- System call will trap into the OS kernel; while library function does not.
  - So, system call is MUCH more slow than library function.
- Library function is the same as the user-defined function. We can replace an existing library function with our own versions, but we can't replace a system call.

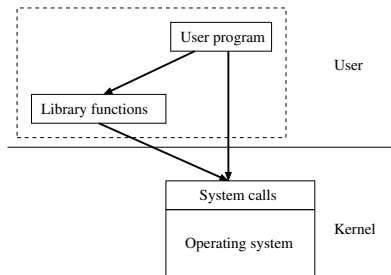


# System call v.s. library function

- System call will trap into the OS kernel; while library function does not.
  - So, system call is MUCH more slow than library function.
- Library function is the same as the user-defined function. We can replace an existing library function with our own versions, but we can't replace a system call.
- A system call in one operating system may become a library function in another operating system and vice versa.

# System call v.s. library function

- System call will trap into the OS kernel; while library function does not.
  - So, system call is MUCH more slow than library function.
- Library function is the same as the user-defined function. We can replace an existing library function with our own versions, but we can't replace a system call.
- A system call in one operating system may become a library function in another operating system and vice versa.



# Questions

- Any questions?



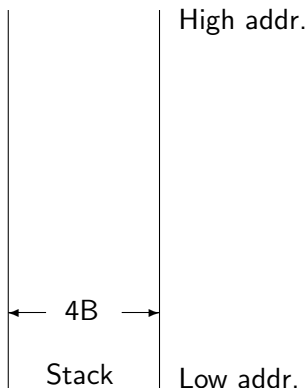
# Function calling convention in C

# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```

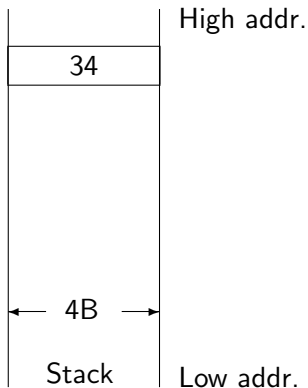
# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



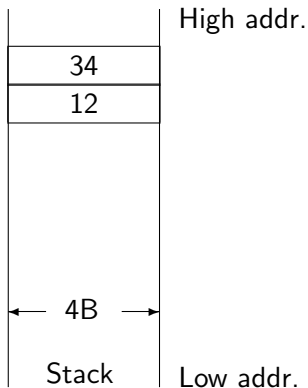
# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

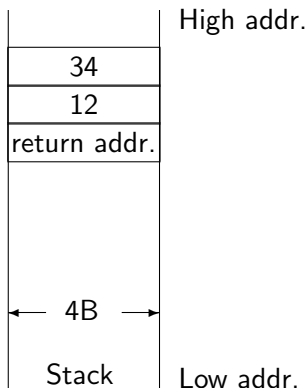
```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```





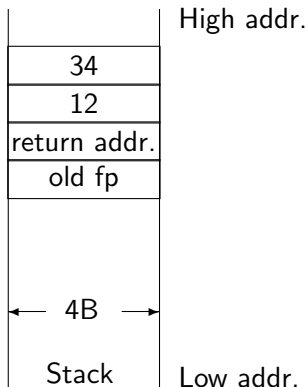
# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



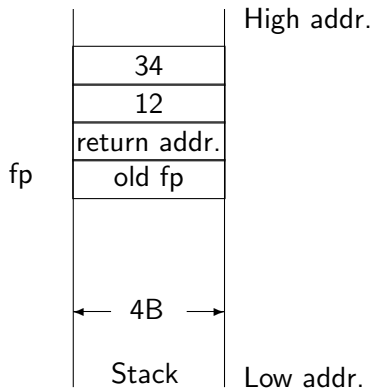
# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



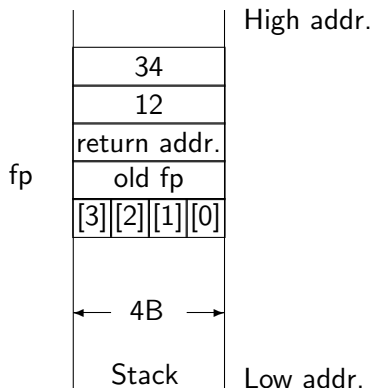
# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



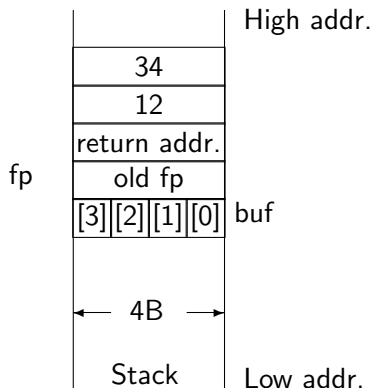
# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

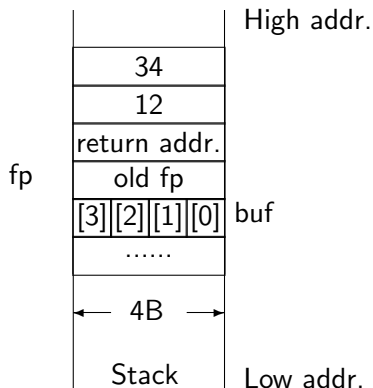
```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

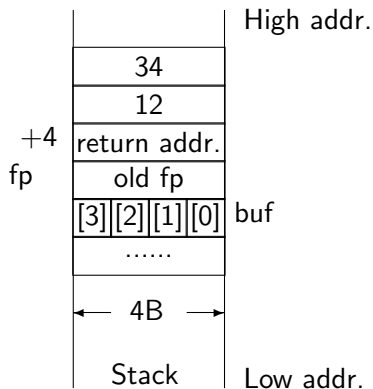
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

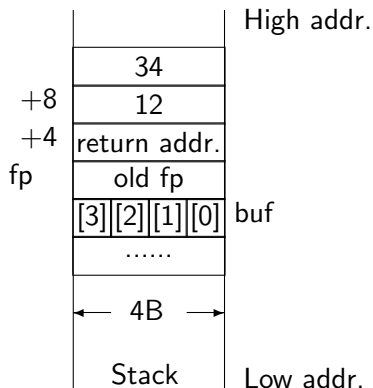
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```

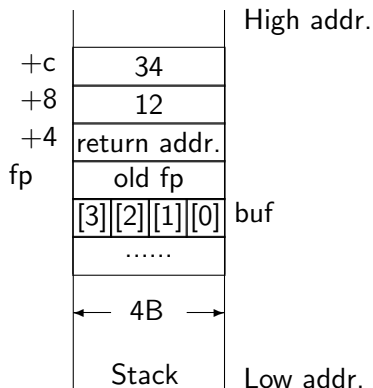




# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

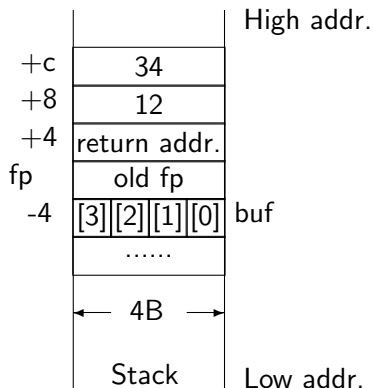
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

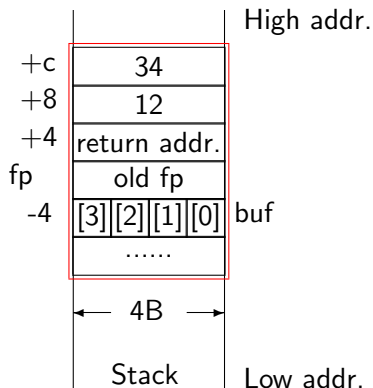
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

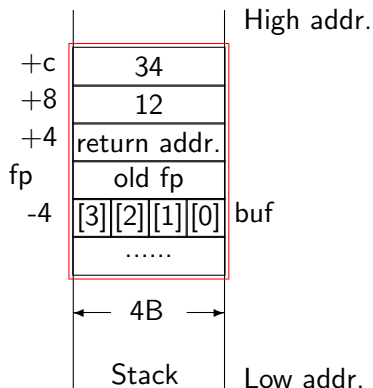
int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```

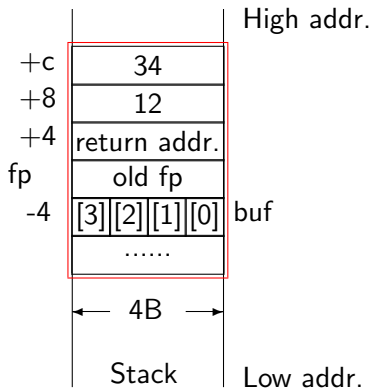


- The area enclosed by red rectangle is called **stack frame**, or simply **frame**.

# Function calling convention in C

```
int foo(int a, int b)
{
    char buf[4];
    return a+b;
}

int main()
{
    int i;
    i = foo(12, 34);
    return i;
}
```



- The area enclosed by red rectangle is called **stack frame**, or simply **frame**.
  - Stack frames are chained into a singly-linked list via **fp** (short for **frame pointer**).

# Stack frame

# Stack frame

- The stack frame (also known as **activation record**) is used to record the information of function calling.

# Stack frame

- The stack frame (also known as **activation record**) is used to record the information of function calling.
  - Parameters



# Stack frame

- The stack frame (also known as **activation record**) is used to record the information of function calling.
  - Parameters
  - Return address

# Stack frame

- The stack frame (also known as **activation record**) is used to record the information of function calling.
  - Parameters
  - Return address
  - Local variables allocation and

# Stack frame

- The stack frame (also known as **activation record**) is used to record the information of function calling.
  - Parameters
  - Return address
  - Local variables allocation and
  - a “prev” pointer to previous stack frame.

# Stack frame

- The stack frame (also known as **activation record**) is used to record the information of function calling.
  - Parameters
  - Return address
  - Local variables allocation and
  - a “prev” pointer to previous stack frame.
- It's stored in the **stack** of the running program.

# Questions

- Any questions?

