

# 操作系统原理

## 第三章：操作系统结构

洪明坚

重庆大学软件学院

February 19, 2016

## 1 Operating system structures

- Simple structure
- Layered structure
- Micro-kernel
- Virtual machine

## 2 Operating system design

- Policy and mechanism
- Implementation

# Outline

## 1 Operating system structures

- Simple structure
- Layered structure
- Micro-kernel
- Virtual machine

## 2 Operating system design

- Policy and mechanism
- Implementation

# System structure

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.
  - Each of these components should be a well-defined portion of the system, with carefully defined inputs, outputs and function.

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.
  - Each of these components should be a well-defined portion of the system, with carefully defined inputs, outputs and function.
- How does the system designer organize these components?



# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.
  - Each of these components should be a well-defined portion of the system, with carefully defined inputs, outputs and function.
- How does the system designer organize these components?
  - Simple structure (or no structure)

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.
  - Each of these components should be a well-defined portion of the system, with carefully defined inputs, outputs and function.
- How does the system designer organize these components?
  - Simple structure (or no structure)
  - Layered structure

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.
  - Each of these components should be a well-defined portion of the system, with carefully defined inputs, outputs and function.
- How does the system designer organize these components?
  - Simple structure (or no structure)
  - Layered structure
  - Micro-kernel

# System structure

- A system as large and complex as a modern operating system must be engineered carefully if it's to function properly and be modified easily.
- A common approach is to partition the task into small components, rather than have one monolithic system.
  - Each of these components should be a well-defined portion of the system, with carefully defined inputs, outputs and function.
- How does the system designer organize these components?
  - Simple structure (or no structure)
  - Layered structure
  - Micro-kernel
  - Virtual machine

# Simple structure

# Simple structure

- Many systems do not have a well-defined structure.

# Simple structure

- Many systems do not have a well-defined structure.
  - They started as small, simple and limited system, and then evolved into a complex one.


# Simple structure

- Many systems do not have a well-defined structure.
  - They started as small, simple and limited system, and then evolved into a complex one.
- Examples: MS-DOS and Unix



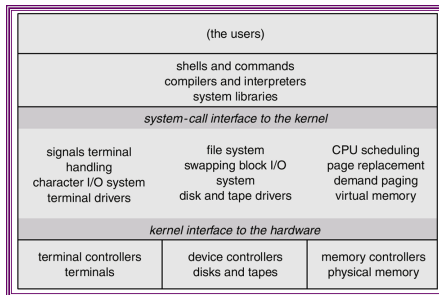

# Simple structure

- Many systems do not have a well-defined structure.
  - They started as small, simple and limited system, and then evolved into a complex one.
- Examples: MS-DOS and Unix



# Simple structure

- Many systems do not have a well-defined structure.
  - They started as small, simple and limited system, and then evolved into a complex one.
- Examples: MS-DOS and Unix



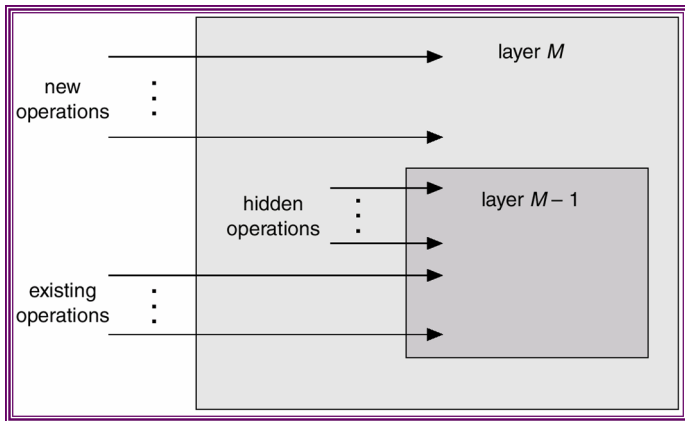
# Layered structure (1/4)

# Layered structure (1/4)

- The operating system is broken up into a number of layers, each built on top lower layers.

# Layered structure (1/4)

- The operating system is broken up into a number of layers, each built on top of lower layers.



# Layered structure (2/4)

## Layered structure (2/4)

- Example 1:

# Layered structure (2/4)

- Example 1:
  - The **THE** operating system by Dijkstra.

Layer	Function
5	The operator
4	User programs
3	Input/output management
2	Operator-process communication
1	Memory and drum management
0	Processor allocation and multiprogramming



# Layered structure (3/4)

# Layered structure (3/4)

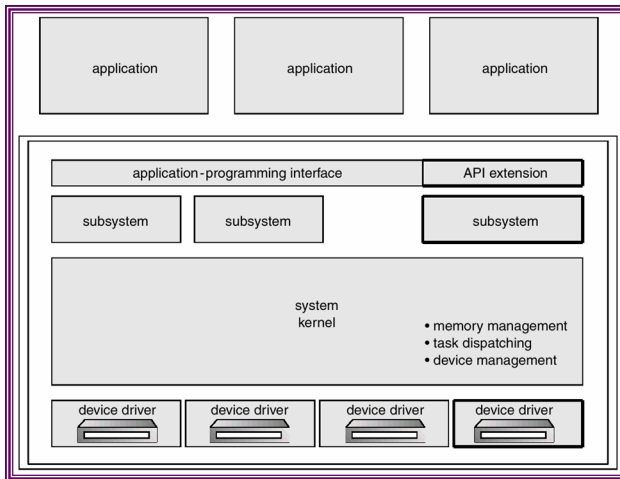
- Example 2:

# Layered structure (3/4)

- Example 2:
  - The IBM OS/2 operating system

# Layered structure (3/4)

- Example 2:
  - The IBM OS/2 operating system



# Layered structure (4/4)

# Layered structure (4/4)

- The major difficulty with the layered structure involves

# Layered structure (4/4)

- The major difficulty with the layered structure involves
  - ① careful definition of the layers and

# Layered structure (4/4)

- The major difficulty with the layered structure involves
  - ① careful definition of the layers and
  - ② less efficient.



# Micro-kernel (1/2)

# Micro-kernel (1/2)

- As the Unix operating system expanded, the kernel became large and difficult to manage.

# Micro-kernel (1/2)

- As the Unix operating system expanded, the kernel became large and difficult to manage.
  - The **micro-kernel** approach structures the operating system by removing all non-essential components from the kernel and implementing them as system- and user-level programs.

# Micro-kernel (1/2)

- As the Unix operating system expanded, the kernel became large and difficult to manage.
  - The **micro-kernel** approach structures the operating system by removing all non-essential components from the kernel and implementing them as system- and user-level programs.
- Which components should remain in the micro-kernel?

# Micro-kernel (1/2)

- As the Unix operating system expanded, the kernel became large and difficult to manage.
  - The **micro-kernel** approach structures the operating system by removing all non-essential components from the kernel and implementing them as system- and user-level programs.
- Which components should remain in the micro-kernel?
  - CPU management

# Micro-kernel (1/2)


- As the Unix operating system expanded, the kernel became large and difficult to manage.
  - The **micro-kernel** approach structures the operating system by removing all non-essential components from the kernel and implementing them as system- and user-level programs.
- Which components should remain in the micro-kernel?
  - CPU management
  - Memory management

# Micro-kernel (1/2)

- As the Unix operating system expanded, the kernel became large and difficult to manage.
  - The **micro-kernel** approach structures the operating system by removing all non-essential components from the kernel and implementing them as system- and user-level programs.
- Which components should remain in the micro-kernel?
  - CPU management
  - Memory management
  - Communication facility

# Micro-kernel (1/2)

- As the Unix operating system expanded, the kernel became large and difficult to manage.
  - The **micro-kernel** approach structures the operating system by removing all non-essential components from the kernel and implementing them as system- and user-level programs.
- Which components should remain in the micro-kernel?
  - CPU management
  - Memory management
  - Communication facility





# Micro-kernel (2/2)

# Micro-kernel (2/2)

- Example 1

# Micro-kernel (2/2)

- Example 1
  - The open source **Mach** from Carnegie Mellon University.

# Micro-kernel (2/2)

- Example 1
  - The open source **Mach** from Carnegie Mellon University.
    - Used as the kernel of the Apple Mac OS X and the DEC Tru64 Unix.

# Micro-kernel (2/2)

- Example 1

- The open source **Mach** from Carnegie Mellon University.
  - Used as the kernel of the Apple Mac OS X and the DEC Tru64 Unix.

- Example 2

# Micro-kernel (2/2)

- Example 1

- The open source **Mach** from Carnegie Mellon University.
  - Used as the kernel of the Apple Mac OS X and the DEC Tru64 Unix.

- Example 2

- QNX real-time operating system by QNX Inc.

# Micro-kernel (2/2)

- Example 1
  - The open source **Mach** from Carnegie Mellon University.
    - Used as the kernel of the Apple Mac OS X and the DEC Tru64 Unix.
- Example 2
  - QNX real-time operating system by QNX Inc.
- Example 3

# Micro-kernel (2/2)

- Example 1
  - The open source **Mach** from Carnegie Mellon University.
    - Used as the kernel of the Apple Mac OS X and the DEC Tru64 Unix.
- Example 2
  - QNX real-time operating system by QNX Inc.
- Example 3
  - Microsoft Windows NT/XP



# Micro-kernel (2/2)

- Example 1


- The open source **Mach** from Carnegie Mellon University.
  - Used as the kernel of the Apple Mac OS X and the DEC Tru64 Unix.

- Example 2

- QNX real-time operating system by QNX Inc.

- Example 3

- Microsoft Windows NT/XP



# Questions

- Any questions?



# Virtual machine (1/4)

# Virtual machine (1/4)

- If we take a step further from micro-kernel,

# Virtual machine (1/4)

- If we take a step further from micro-kernel,
  - The low-level real hardware is “cloned” into several identical **virtual machines**.

# Virtual machine (1/4)

- If we take a step further from micro-kernel,
  - The low-level real hardware is “cloned” into several identical **virtual machines**.
  - That is, virtual machine provides an interface identical to the underlying bare hardware.

# Virtual machine (1/4)

- If we take a step further from micro-kernel,
  - The low-level real hardware is “cloned” into several identical **virtual machines**.
  - That is, virtual machine provides an interface identical to the underlying bare hardware.
- Then, the operating system functionality is built on top of the virtual machines.

# Virtual machine (2/4)

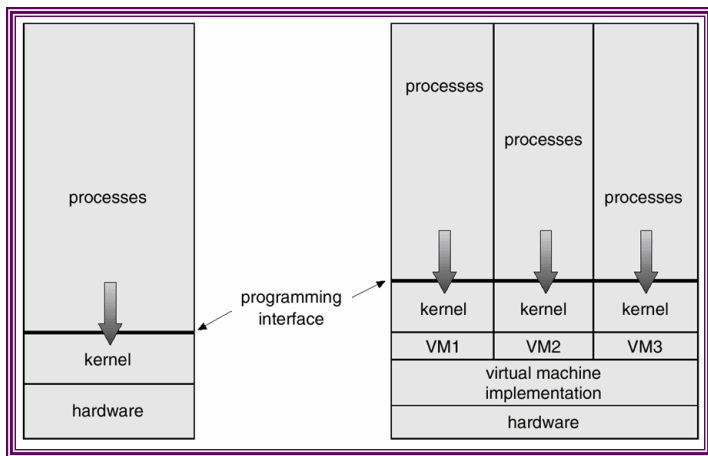


# Virtual machine (2/4)

- Virtual machine structure


# Virtual machine (2/4)

- Virtual machine structure



# Virtual machine (2/4)

- Virtual machine structure



# Virtual machine (3/4)

# Virtual machine (3/4)

- Examples

# Virtual machine (3/4)

- Examples
  - IBM VM/370

# Virtual machine (3/4)

- Examples
  - IBM VM/370
  - VMware

# Virtual machine (3/4)

- Examples
  - IBM VM/370
  - VMware
  - Microsoft Virtual PC



# Virtual machine (3/4)

- Examples

- IBM VM/370
- VMware
- Microsoft Virtual PC
- Sun microsystem's Java Virtual Machine (JVM).

# Virtual machine (4/4)

# Virtual machine (4/4)

- Advantage and disadvantage of virtual machine

# Virtual machine (4/4)

- Advantage and disadvantage of virtual machine
  - The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines.

# Virtual machine (4/4)

- Advantage and disadvantage of virtual machine
  - The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines.
    - This isolation, however, permits no direct sharing of resources

# Virtual machine (4/4)

- Advantage and disadvantage of virtual machine
  - The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines.
    - This isolation, however, permits no direct sharing of resources
  - A virtual-machine system is a perfect vehicle for operating-systems research and development.

# Virtual machine (4/4)

- Advantage and disadvantage of virtual machine
  - The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines.
    - This isolation, however, permits no direct sharing of resources
  - A virtual-machine system is a perfect vehicle for operating-systems research and development.
  - The virtual machine concept is difficult to implement due to the effort required to provide an **exact** duplicate to the underlying machine.

# Questions

- Any questions?





# Outline

## 1 Operating system structures

- Simple structure
- Layered structure
- Micro-kernel
- Virtual machine

## 2 Operating system design

- Policy and mechanism
- Implementation

# Operating System design

# Operating System design

- Design goals

# Operating System design

- Design goals

**User goals** operating system should be convenient to use, easy to learn, reliable, safe, and fast.

- Design goals

**User goals** operating system should be convenient to use, easy to learn, reliable, safe, and fast.

**System goals** operating system should be easy to design, implement, and maintain, as well as flexible, reliable, error-free, and efficient.

# Policy and mechanism

- **Policies**

- **Policies - what to do**



- **Policies - what to do**

- For example, users should not be able to read other users' files.

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms**

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

- It allows maximum flexibility if policy decisions are to be changed later.

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

- It allows maximum flexibility if policy decisions are to be changed later.

- Two extremes:

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

- It allows maximum flexibility if policy decisions are to be changed later.

- Two extremes:

- Micro-kernel



# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

- It allows maximum flexibility if policy decisions are to be changed later.

- Two extremes:

- Micro-kernel - all mechanisms, almost no policy

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

- It allows maximum flexibility if policy decisions are to be changed later.

- Two extremes:

- Micro-kernel - all mechanisms, almost no policy
- The Apple Macintosh

# Policy and mechanism

- **Policies - what to do**

- For example, users should not be able to read other users' files.

- **Mechanisms - how to do**

- For example, file permissions are checked on *open* system call.

- The **separation** of policy from mechanism is a very important principle

- It allows maximum flexibility if policy decisions are to be changed later.

- Two extremes:

- Micro-kernel - all mechanisms, almost no policy
- The Apple Macintosh - policy and mechanism are bound together

# The operating system implementation

# The operating system implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages.

# The operating system implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages.
- Code written in a high-level language:

# The operating system implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages.
- Code written in a high-level language:
  - can be written faster.

# The operating system implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages.
- Code written in a high-level language:
  - can be written faster.
  - is more compact.



# The operating system implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages.
- Code written in a high-level language:
  - can be written faster.
  - is more compact.
  - is easier to understand and debug.

# The operating system implementation

- Traditionally written in assembly language, operating systems can now be written in higher-level languages.
- Code written in a high-level language:
  - can be written faster.
  - is more compact.
  - is easier to understand and debug.
  - is easier to port to some other hardwares.

# Questions

- Any questions?

