

操作系统原理

第七章：进程同步

洪明坚

重庆大学软件学院

February 19, 2016

目录

- 1 Race condition
- 2 Possible solutions to race condition
 - Software solutions
 - Hardware solutions
 - Conclusion
- 3 Semaphore
 - What's the semaphore?
 - Programming interfaces
 - Classic problems of synchronization
 - Binary semaphore
 - Deadlock and starvation
- 4 Monitor
 - Condition variables
 - Monitor solution to the dining-philosopher problem
 - Language support
- 5 Relationships of semaphore and monitor

Outline

- 1 Race condition
- 2 Possible solutions to race condition
 - Software solutions
 - Hardware solutions
 - Conclusion
- 3 Semaphore
 - What's the semaphore?
 - Programming interfaces
 - Classic problems of synchronization
 - Binary semaphore
 - Deadlock and starvation
- 4 Monitor
 - Condition variables
 - Monitor solution to the dining-philosopher problem
 - Language support
- 5 Relationships of semaphore and monitor

Background

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.
 - Multiple threads within a process may share a piece of memory by using global variables.

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.
 - Multiple threads within a process may share a piece of memory by using global variables.
- Concurrent access to shared data may result in data inconsistency.

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.
 - Multiple threads within a process may share a piece of memory by using global variables.
- Concurrent access to shared data may result in data inconsistency.
 - Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.
 - Multiple threads within a process may share a piece of memory by using global variables.
- Concurrent access to shared data may result in data inconsistency.
 - Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Solution to *producer and consumer problem* (Chapter 4) allows at most $(BSIZE - 1)$ items in buffer at the same time.

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.
 - Multiple threads within a process may share a piece of memory by using global variables.
- Concurrent access to shared data may result in data inconsistency.
 - Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Solution to *producer and consumer problem* (Chapter 4) allows at most $(BSIZE - 1)$ items in buffer at the same time.
 - A solution, where all $BSIZE$ buffers are used is NOT simple.

Background

- Cooperating processes may share a block of memory via IPC facilities provided by the kernel.
 - Multiple threads within a process may share a piece of memory by using global variables.
- Concurrent access to shared data may result in data inconsistency.
 - Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Solution to *producer and consumer problem* (Chapter 4) allows at most $(BSIZE - 1)$ items in buffer at the same time.
 - A solution, where all $BSIZE$ buffers are used is NOT simple.
 - Suppose that we modify the code by adding a variable *counter*.

Producer and consumer problem

Producer and consumer problem

- Shared-memory bounded-buffer

Producer and consumer problem

- Shared-memory bounded-buffer

```
/* Shared variables */
#define BSIZE 10
struct item {
    ....
} buffer[BSIZE];
int in = 0, out = 0, counter = 0;

/*-----*/
/*The producer loop*/
while(1) {
    /*produce an item*/
    while(counter == BSIZE)
        /*do nothing*/;
    buffer[in] = itemProduced;
    in = (in + 1) % BSIZE;
    counter++;
}

/*The consumer loop*/
while(1) {
    while(counter == 0)
        /*do nothing*/;
    itemConsumed = buffer[out];
    out = (out + 1) % BSIZE;
    counter--;
    /*consume the item*/
}
```

Problem of above solution (1/2)

Problem of above solution (1/2)

- When producer and consumer processes are executed *concurrently*, they may not function correctly.

Problem of above solution (1/2)

- When producer and consumer processes are executed *concurrently*, they may not function correctly.
 - We can show that the value of *counter* may be incorrect as follows.

Problem of above solution (1/2)

- When producer and consumer processes are executed *concurrently*, they may not function correctly.
 - We can show that the value of *counter* may be incorrect as follows.
- Note that the “*counter++*” and “*counter--*” may be implemented in machine language as

Problem of above solution (1/2)

- When producer and consumer processes are executed *concurrently*, they may not function correctly.
 - We can show that the value of *counter* may be incorrect as follows.
- Note that the “*counter++*” and “*counter--*” may be implemented in machine language as

```
; counter++
register1 = counter      ; load the counter to a register
register1 = register1 + 1 ; add
counter = register1      ; write back

; counter--
register2 = counter      ; load the counter to a register
register2 = register2 - 1 ; subtract
counter = register2      ; write back
```

Problem of above solution (2/2)

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1= <i>counter</i>		{register1=5}

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1= <i>counter</i>		{register1=5}
T_1	register1=register1+1		{register1=6}

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1=counter		{register1=5}
T_1	register1=register1+1		{register1=6}
T_2		register2=counter	{register2=5}

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1=counter		{register1=5}
T_1	register1=register1+1		{register1=6}
T_2		register2=counter	{register2=5}
T_3		register2=register2-1	{register2=4}

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1=counter		{register1=5}
T_1	register1=register1+1		{register1=6}
T_2		register2=counter	{register2=5}
T_3		register2=register2-1	{register2=4}
T_4	counter=register1		{counter=6}

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1=counter		{register1=5}
T_1	register1=register1+1		{register1=6}
T_2		register2=counter	{register2=5}
T_3		register2=register2-1	{register2=4}
T_4	counter=register1		{counter=6}
T_5		counter=register2	{counter=4}

Problem of above solution (2/2)

- It's important to note that the above two instruction sequences can be *interleaved* because of interrupts or scheduling.
- For example, assume *counter* is initially 5,

Time	Producer	Consumer	result
T_0	register1=counter		{register1=5}
T_1	register1=register1+1		{register1=6}
T_2		register2=counter	{register2=5}
T_3		register2=register2-1	{register2=4}
T_4	counter=register1		{counter=6}
T_5		counter=register2	{counter=4}

- The value of *counter* may be either 4 or 6, while the correct result should be 5.

Questions

- Any questions?



Race condition (1/2)

Race condition (1/2)

- *Race condition is a situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.*

Race condition (1/2)

- *Race condition is a situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.*
- How do we avoid race condition?

Race condition (1/2)

- *Race condition is a situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.*
- How do we avoid race condition?
 - Find some way to prohibit more than one process from reading and writing the shared data concurrently.

Race condition (1/2)

- *Race condition is a situation where several processes access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.*
- How do we avoid race condition?
 - Find some way to prohibit more than one process from reading and writing the shared data concurrently.
 - That is, the access must be *serialized* even if the processes attempt concurrent access.

Race condition (2/2)

Race condition (2/2)

- **Critical section**

Race condition (2/2)

- **Critical section**

- a piece of code where the shared resource is accessed.

Race condition (2/2)

- **Critical section**

- a piece of code where the shared resource is accessed.
- A solution of race condition must satisfy the following 4 requirements:

Race condition (2/2)

- **Critical section**

- a piece of code where the shared resource is accessed.

- A solution of race condition must satisfy the following 4 requirements:

- 1 **Mutual exclusion:** No two process may be simultaneously inside their critical section;

Race condition (2/2)

- **Critical section**

- a piece of code where the shared resource is accessed.

- A solution of race condition must satisfy the following 4 requirements:

- ① **Mutual exclusion:** No two process may be simultaneously inside their critical section;
- ② **Progress:** No process running outside its critical section may block other processes trying to enter its critical section;

Race condition (2/2)

- **Critical section**

- a piece of code where the shared resource is accessed.

- A solution of race condition must satisfy the following 4 requirements:

- 1 **Mutual exclusion:** No two process may be simultaneously inside their critical section;
- 2 **Progress:** No process running outside its critical section may block other processes trying to enter its critical section;
- 3 **Bounded waiting:** No process should have to wait forever to enter its critical section;

Race condition (2/2)

- **Critical section**

- a piece of code where the shared resource is accessed.

- A solution of race condition must satisfy the following 4 requirements:

- ❶ **Mutual exclusion:** No two process may be simultaneously inside their critical section;
- ❷ **Progress:** No process running outside its critical section may block other processes trying to enter its critical section;
- ❸ **Bounded waiting:** No process should have to wait forever to enter its critical section;
- ❹ **Speed:** No assumptions may be made about speeds or the number of the CPUs.

Critical section

Critical section

- A protocol must be designed to be used by the processes to enter and leave critical section.

Critical section

- A protocol must be designed to be used by the processes to enter and leave critical section.
 - Each process must request permission to enter its critical section.

Critical section

- A protocol must be designed to be used by the processes to enter and leave critical section.
 - Each process must request permission to enter its critical section.

```
do {  
    <entry section>  
  
    CRITICAL SECTION  
  
    <exit section>  
} while (1);
```


Critical section

- A protocol must be designed to be used by the processes to enter and leave critical section.
 - Each process must request permission to enter its critical section.

```
do {  
    <entry section>  
  
    CRITICAL SECTION  
  
    <exit section>  
  
} while (1);
```

- In the following slides, we will try to construct several “entry section” and “exit section” to solve the race condition.

Questions

- Any questions?



Outline

- 1 Race condition
- 2 Possible solutions to race condition
 - Software solutions
 - Hardware solutions
 - Conclusion
- 3 Semaphore
 - What's the semaphore?
 - Programming interfaces
 - Classic problems of synchronization
 - Binary semaphore
 - Deadlock and starvation
- 4 Monitor
 - Condition variables
 - Monitor solution to the dining-philosopher problem
 - Language support
- 5 Relationships of semaphore and monitor

First try

First try

- Assume there are only two cooperating processes: P_i and P_j , where $i+j=1$.

First try

- Assume there are only two cooperating processes: P_i and P_j , where $i+j=1$.

```
int turn = 0; // or 1
do {
    while(turn != i); // entry section

    CRITICAL SECTION

    turn = j;          // exit section
} while (1);
```

- Does it satisfy all 4 requirements for a solution?

First try

- Assume there are only two cooperating processes: P_i and P_j , where $i+j=1$.

```
int turn = 0; // or 1
do {
    while(turn != i); // entry section

    CRITICAL SECTION

    turn = j;          // exit section
} while (1);
```

- Does it satisfy all 4 requirements for a solution?
 - No, it breaks the **progress** requirement.

Second try

Second try

```
bool flag[2] = {false, false};  
do {  
    // entry section  
    flag[i] = true;  
    while(flag[j]);  
  
    CRITICAL SECTION  
  
    flag[i] = false; // exit section  
} while (1);
```

- Does it satisfy all 4 requirements for a solution?

Second try

```
bool flag[2] = {false, false};  
do {  
    // entry section  
    flag[i] = true;  
    while(flag[j]);  
  
    CRITICAL SECTION  
  
    flag[i] = false; // exit section  
} while (1);
```

- Does it satisfy all 4 requirements for a solution?
 - No, it breaks the **progress** requirement too.

Third try

Third try

```
int turn = 0; // or 1
bool flag[2] = {false, false};
do {
    // entry section
    flag[i] = true;
    turn = j;
    while(flag[j] && turn == j);

    CRITICAL SECTION

    flag[i] = false; // exit section
} while (1);
```

- Does it satisfy all 4 requirements for a solution?

Third try

```
int turn = 0; // or 1
bool flag[2] = {false, false};
do {
    // entry section
    flag[i] = true;
    turn = j;
    while(flag[j] && turn == j);

    CRITICAL SECTION

    flag[i] = false; // exit section
} while (1);
```

- Does it satisfy all 4 requirements for a solution?
 - Yes, it's a correct solution and is known as *Peterson's algorithm*.

Bakery algorithm

Bakery algorithm

- Peterson's algorithm solves the problem for two processes, while *Bakery algorithm* solves it for multiple processes.

Bakery algorithm

- Peterson's algorithm solves the problem for two processes, while *Bakery algorithm* solves it for multiple processes.

```
bool choosing[n] = {false, ..., false};
int number[n] = {0, ..., 0};
do {
    // entry section
    choosing[i] = true;
    number[i] = max(number[0], ..., number[n-1])+1;
    choosing[i] = false;
    for(j=0; j < n; j++) {
        while(choosing[j]);
        while((number[j]!=0)&&(number[j],j)<(number[i],i));
    }
}
```

CRITICAL SECTION

```
number[i] = 0; // exit section
```

```
} while (1);
```


Questions

- Any questions?



Synchronization hardware

Synchronization hardware

- Several low-level hardware features may be used to solve the race condition.

Synchronization hardware

- Several low-level hardware features may be used to solve the race condition.
 - ① Disabling interrupts

Synchronization hardware

- Several low-level hardware features may be used to solve the race condition.
 - ① Disabling interrupts
 - ② Special instructions

Synchronization hardware

- Several low-level hardware features may be used to solve the race condition.
 - ① Disabling interrupts
 - ② Special instructions
 - TSL (Test and Set Lock)

Synchronization hardware

- Several low-level hardware features may be used to solve the race condition.
 - ① Disabling interrupts
 - ② Special instructions
 - TSL (Test and Set Lock)
 - SWAP

Disabling interrupts

Disabling interrupts

- The CPU is only switched from process to process as a result of clock or other interrupts.

Disabling interrupts

- The CPU is only switched from process to process as a result of clock or other interrupts.
 - Once a process has disabled interrupts, it can access the shared memory without fear that any other process will intervene.

Disabling interrupts

- The CPU is only switched from process to process as a result of clock or other interrupts.
 - Once a process has disabled interrupts, it can access the shared memory without fear that any other process will intervene.
- Disadvantages

Disabling interrupts

- The CPU is only switched from process to process as a result of clock or other interrupts.
 - Once a process has disabled interrupts, it can access the shared memory without fear that any other process will intervene.
- Disadvantages
 - ① User processes should NOT be able to disable the interrupts;

Disabling interrupts

- The CPU is only switched from process to process as a result of clock or other interrupts.
 - Once a process has disabled interrupts, it can access the shared memory without fear that any other process will intervene.
- Disadvantages
 - ① User processes should NOT be able to disable the interrupts;
 - ② It's not feasible in a multiprocessor system.

TSL and SWAP (1/3)

TSL and SWAP (1/3)

- The TSL and SWAP instructions have the following functionalities, respectively:

TSL and SWAP (1/3)

- The TSL and SWAP instructions have the following functionalities, respectively:

```
bool TSL(bool &target)
{
    bool rv = target;
    target = true;
    return rv;
}
```

```
void SWAP(bool &a, bool &b)
{
    bool temp = a;
    a = b;
    b = temp;
}
```


TSL and SWAP (1/3)

- The TSL and SWAP instructions have the following functionalities, respectively:

```
bool TSL(bool &target)
{
    bool rv = target;
    target = true;
    return rv;
}
```

```
void SWAP(bool &a, bool &b)
{
    bool temp = a;
    a = b;
    b = temp;
}
```

- Bear in mind that TSL and SWAP are executed *atomically*, that is, as one uninterruptible unit.

TSL and SWAP (2/3)

TSL and SWAP (2/3)

- Use TSL or SWAP to TRY to solve race condition

TSL and SWAP (2/3)

- Use TSL or SWAP to TRY to solve race condition

```
bool lock = false;  
do {  
    // entry section  
    while(TSL(lock))  
        ;
```

CRITICAL SECTION

```
    // exit section  
    lock = false;
```

```
} while (1);
```

```
bool lock=false;  
do {  
    // entry section  
    bool key = true;  
    while(key == true)  
        SWAP(lock , key);
```

CRITICAL SECTION

```
    // exit section  
    lock = false;
```

```
} while (1);
```

TSL and SWAP (2/3)

- Use TSL or SWAP to TRY to solve race condition

```
bool lock = false;  
do {  
    // entry section  
    while(TSL(lock))  
        ;
```

CRITICAL SECTION

```
// exit section  
lock = false;
```

```
} while (1);
```

```
bool lock=false;  
do {  
    // entry section  
    bool key = true;  
    while(key == true)  
        SWAP(lock , key);
```

CRITICAL SECTION

```
// exit section  
lock = false;
```

```
} while (1);
```

- But these two algorithms do NOT satisfy the **bounded-waiting** requirement.

TSL and SWAP (3/3)

TSL and SWAP (3/3)

- A correct solution using TSL

TSL and SWAP (3/3)

- A correct solution using TSL

```
bool lock = false , waiting[n] = {false , ... , false};
do {
    // entry section
    waiting[i] = true;  bool key = true;
    while(waiting[i] && key)
        key = TSL(lock);
    waiting[i] = false;
```

CRITICAL SECTION

```
    // exit section
    int j = (i + 1) % n;
    while((j != i) && !waiting[j])
        j = (j+1) % n;
    if(j == i)    lock = false;
    else        waiting[j] = false;

} while (1);
```


Conclusion

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.
 - While a process is in its critical section, any other process tries to enter its critical section must loop continuously in the entry code.

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.
 - While a process is in its critical section, any other process tries to enter its critical section must loop continuously in the entry code.
 - Busy waiting wastes CPU cycles that some other process might be able to use productively.

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.
 - While a process is in its critical section, any other process tries to enter its critical section must loop continuously in the entry code.
 - Busy waiting wastes CPU cycles that some other process might be able to use productively.
- This type of solution is also called a *spinlock*.

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.
 - While a process is in its critical section, any other process tries to enter its critical section must loop continuously in the entry code.
 - Busy waiting wastes CPU cycles that some other process might be able to use productively.
- This type of solution is also called a *spinlock*.
 - Because the process “spins” while waiting on a lock.

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.
 - While a process is in its critical section, any other process tries to enter its critical section must loop continuously in the entry code.
 - Busy waiting wastes CPU cycles that some other process might be able to use productively.
- This type of solution is also called a *spinlock*.
 - Because the process “spins” while waiting on a lock.
 - The spinlocks are **only** useful in multiprocessor systems.

Conclusion

- The above solutions have a common disadvantage: *busy waiting*.
 - While a process is in its critical section, any other process tries to enter its critical section must loop continuously in the entry code.
 - Busy waiting wastes CPU cycles that some other process might be able to use productively.
- This type of solution is also called a *spinlock*.
 - Because the process “spins” while waiting on a lock.
 - The spinlocks are **only** useful in multiprocessor systems.
 - Because no context switch is required if the locks are expected to be held for short times.

Questions

- Any questions?



Outline

- 1 Race condition
- 2 Possible solutions to race condition
 - Software solutions
 - Hardware solutions
 - Conclusion
- 3 Semaphore
 - What's the semaphore?
 - Programming interfaces
 - Classic problems of synchronization
 - Binary semaphore
 - Deadlock and starvation
- 4 Monitor
 - Condition variables
 - Monitor solution to the dining-philosopher problem
 - Language support
- 5 Relationships of semaphore and monitor

Semaphore

- The above solutions to race condition are not easy to generalize to more complex problems.

Semaphore

- The above solutions to race condition are not easy to generalize to more complex problems.
- To overcome this difficulty, we can use a synchronization tool called *semaphore*.

Semaphore

- The above solutions to race condition are not easy to generalize to more complex problems.
- To overcome this difficulty, we can use a synchronization tool called *semaphore*.
 - It's invented by *Edsger Dijkstra* and first used in the *THE* operating system.

Semaphore

- The above solutions to race condition are not easy to generalize to more complex problems.
- To overcome this difficulty, we can use a synchronization tool called *semaphore*.
 - It's invented by *Edsger Dijkstra* and first used in the *THE* operating system.
- What's the semaphore?

Semaphore

- The above solutions to race condition are not easy to generalize to more complex problems.
- To overcome this difficulty, we can use a synchronization tool called *semaphore*.
 - It's invented by *Edsger Dijkstra* and first used in the *THE* operating system.
- What's the semaphore?
 - A semaphore S is an integer variable that, apart from initialization, is accessed **only** through two standard **atomic** operations: $P(down)$ and $V(up)$.

Implementation (1/2)

Implementation (1/2)

```
typedef struct {  
    int value;  
  
    /*processes blocked by this semaphore*/  
    struct PCB *L;  
} semaphore;
```

```
void P(semaphore *S) {  
    S->value--;  
    if(S->value < 0) {  
        Add (curproc) to S->L;  
  
        (curproc)->state=WAITING;  
        scheduler();  
    }  
}
```

```
void V(semaphore *S) {  
    S->value++;  
    if(S->value <= 0) {  
        Remove a proc (A) from S->L;  
  
        (A)->state=READY; //Wakeup (A)  
    }  
}
```

Implementation (2/2)

Implementation (2/2)

- The magnitude of the “value” is the number of resources available (>0) or of processes waiting on that semaphore (<0).

Implementation (2/2)

- The magnitude of the “value” is the number of resources available (>0) or of processes waiting on that semaphore (<0).
- The P and V must be executed **atomically**.

Implementation (2/2)

- The magnitude of the “value” is the number of resources available (>0) or of processes waiting on that semaphore (<0).
- The P and V must be executed **atomically**.
 - In either of two ways:

Implementation (2/2)

- The magnitude of the “value” is the number of resources available (>0) or of processes waiting on that semaphore (<0).
- The P and V must be executed **atomically**.
 - In either of two ways:
 - ① Disable interrupts in uni-processor systems or

Implementation (2/2)

- The magnitude of the “value” is the number of resources available (>0) or of processes waiting on that semaphore (<0).
- The P and V must be executed **atomically**.
 - In either of two ways:
 - ① Disable interrupts in uni-processor systems or
 - ② Spinlocks in multi-processor systems.

Application programming interface

Application programming interface

- Win32

Application programming interface

- Win32
 - CreateSemaphore/CloseHandle

Application programming interface

- Win32
 - CreateSemaphore/CloseHandle
 - WaitForSingleObject/ReleaseSemaphore

Application programming interface

- Win32
 - CreateSemaphore/CloseHandle
 - WaitForSingleObject/ReleaseSemaphore
- POSIX

Application programming interface

- Win32
 - CreateSemaphore/CloseHandle
 - WaitForSingleObject/ReleaseSemaphore
- POSIX
 - sem_init/sem_destroy

Application programming interface

- Win32
 - CreateSemaphore/CloseHandle
 - WaitForSingleObject/ReleaseSemaphore
- POSIX
 - sem_init/sem_destroy
 - sem_wait/sem_post

Questions

- Any questions?



Classic problems of synchronization

Classic problems of synchronization

- The producer and consumer problem

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem
 - A data object is to be shared among several concurrent processes.

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem
 - A data object is to be shared among several concurrent processes.
 - Some of them (*readers*) may want **only** to read the data object,

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem
 - A data object is to be shared among several concurrent processes.
 - Some of them (*readers*) may want **only** to read the data object, while others (*writers*) may want to update the data object.

Classic problems of synchronization


- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem
 - A data object is to be shared among several concurrent processes.
 - Some of them (*readers*) may want **only** to read the data object, while others (*writers*) may want to update the data object.
 - Multiple readers may read the data object simultaneously, but the writers must have exclusive access to the data object.

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem
 - A data object is to be shared among several concurrent processes.
 - Some of them (*readers*) may want **only** to read the data object, while others (*writers*) may want to update the data object.
 - Multiple readers may read the data object simultaneously, but the writers must have exclusive access to the data object.
- The dining-philosophers problem

Classic problems of synchronization

- The producer and consumer problem
 - Explained in the chapter “Process Management”.
- The readers-writers problem
 - A data object is to be shared among several concurrent processes.
 - Some of them (*readers*) may want **only** to read the data object, while others (*writers*) may want to update the data object.
 - Multiple readers may read the data object simultaneously, but the writers must have exclusive access to the data object.
- The dining-philosophers problem



The producer and consumer problem (1/2)

The producer and consumer problem (1/2)

- A semaphore *mutex* is used to protect the critical section when accessing the buffer.

The producer and consumer problem (1/2)

- A semaphore *mutex* is used to protect the critical section when accessing the buffer.
 - It's initialized to the value 1.

The producer and consumer problem (1/2)

- A semaphore *mutex* is used to protect the critical section when accessing the buffer.
 - It's initialized to the value 1.
- The semaphores *empty* and *full* synchronize the producer and consumer.

The producer and consumer problem (1/2)

- A semaphore *mutex* is used to protect the critical section when accessing the buffer.
 - It's initialized to the value 1.
- The semaphores *empty* and *full* synchronize the producer and consumer.
 - *empty* is initialized to the value *BSIZE*;

The producer and consumer problem (1/2)

- A semaphore *mutex* is used to protect the critical section when accessing the buffer.
 - It's initialized to the value 1.
- The semaphores *empty* and *full* synchronize the producer and consumer.
 - *empty* is initialized to the value *BSIZE*;
 - *full* is initialized to the value 0.

The producer and consumer problem (2/2)

The producer and consumer problem (2/2)

```
/*The producer loop*/  
do {  
    ...  
    produce an item;  
    ...  
    P(&empty);  
    P(&mutex);  
    ...  
    add the item to buffer;  
    ...  
    V(&mutex);  
    V(&full);  
} while(1);
```

The producer and consumer problem (2/2)

```
/*The producer loop*/
do {
    ...
    produce an item;
    ...
    P(&empty);
    P(&mutex);
    ...
    add the item to buffer;
    ...
    V(&mutex);
    V(&full);
} while(1);
```

```
/*The consumer loop*/
do {
    P(&full);
    P(&mutex);
    ...
    remove an item from buffer;
    ...
    V(&mutex);
    V(&empty);
    ...
    consume the item;
    ...
} while(1);
```

Questions

- Any questions?



The readers-writers problem (1/2)

The readers-writers problem (1/2)

- A semaphore *wrt* is used to protect the shared data object.

The readers-writers problem (1/2)

- A semaphore *wrt* is used to protect the shared data object.
 - It's initialized to the value 1.

The readers-writers problem (1/2)

- A semaphore *wrt* is used to protect the shared data object.
 - It's initialized to the value 1.
- An integer *readcount* count the number of readers which is busy reading.

The readers-writers problem (1/2)

- A semaphore *wrt* is used to protect the shared data object.
 - It's initialized to the value 1.
- An integer *readcount* count the number of readers which is busy reading.
 - *readcount* is initialized to the value 0;

The readers-writers problem (1/2)

- A semaphore *wrt* is used to protect the shared data object.
 - It's initialized to the value 1.
- An integer *readcount* count the number of readers which is busy reading.
 - *readcount* is initialized to the value 0;
- Another semaphore *mutex* is used to protect the *readcount*.

The readers-writers problem (1/2)

- A semaphore *wrt* is used to protect the shared data object.
 - It's initialized to the value 1.
- An integer *readcount* count the number of readers which is busy reading.
 - *readcount* is initialized to the value 0;
- Another semaphore *mutex* is used to protect the *readcount*.
 - *mutex* is initialized to the value 1.

The readers-writers problem (2/2)

The readers-writers problem (2/2)

```
/*The writer loop*/  
do {  
    P(&wrt);  
    ...  
    writing is performed  
    ...  
    V(&wrt);  
} while(1);
```

The readers-writers problem (2/2)

```
/*The writer loop*/
do {
    P(&wrt);
    ...
    writing is performed
    ...
    V(&wrt);
} while(1);
```

```
/*The reader loop*/
do {
    P(&mutex);
    readcount++;
    if(readcount == 1)
        P(&wrt);
    V(&mutex);
    ...
    reading is performed
    ...
    P(&mutex);
    readcount--;
    if(readcount == 0)
        V(&wrt);
    V(&mutex);
} while(1);
```

Questions

- Any questions?



The dining-philosopher problem (1/3)

The dining-philosopher problem (1/3)

```
semaphore chopstick[5] = {1, ..., 1};  
/* philosopher i */  
do {  
    ...  
    thinking  
    ...  
    P(&chopstick[i]);  
    P(&chopstick[(i + 1) % 5]);  
    ...  
    eating  
    ...  
    V(&chopstick[(i + 1) % 5]);  
    V(&chopstick[i]);  
} while(1);
```

The dining-philosopher problem (1/3)

```
semaphore chopstick[5] = {1, ..., 1};  
/* philosopher i */  
do {  
    ...  
    thinking  
    ...  
    P(&chopstick[i]);  
    P(&chopstick[(i + 1) % 5]);  
    ...  
    eating  
    ...  
    V(&chopstick[(i + 1) % 5]);  
    V(&chopstick[i]);  
} while(1);
```

- Is this a correct solution?

The dining-philosopher problem (1/3)

```
semaphore chopstick[5] = {1, ..., 1};  
/* philosopher i */  
do {  
    ...  
    thinking  
    ...  
    P(&chopstick[i]);  
    P(&chopstick[(i + 1) % 5]);  
    ...  
    eating  
    ...  
    V(&chopstick[(i + 1) % 5]);  
    V(&chopstick[i]);  
} while(1);
```

- Is this a correct solution?
 - No.

The dining-philosophers problem (2/3)

The dining-philosophers problem (2/3)

```
#define N          5
#define LEFT      ((i-1) % N)
#define RIGHT     ((i+1) % N)

enum {THINKING, HUNGRY, EATING} state[N];
semaphore mutex = 1; /*used to protect the 'state'*/

semaphore s[N] = {0, ..., 0};

void philosopher(int i)
{
    while(1) {
        think();
        take_forks(i);
        eat();
        put_forks(i);
    }
}
```

The dining-philosophers problem (3/3)

The dining-philosophers problem (3/3)

```
void test(int i)
{
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        V(&s[i]);
    }
}
```

The dining-philosophers problem (3/3)

```
void test(int i)
{
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        V(&s[i]);
    }
}
```

```
void take_forks(int i)
{
    P(&mutex);
    state[i] = HUNGRY;
    test(i);
    V(&mutex);
    P(&s[i]);
}
```

```
void put_forks(int i)
{
    P(&mutex);
    state[i] = THINKING;
    test(LEFT);
    test(RIGHT);
    V(&mutex);
}
```


Questions

- Any questions?



Binary semaphore

Binary semaphore

- The semaphore construct described in the previous slides is commonly known as a *counting semaphore*.

Binary semaphore

- The semaphore construct described in the previous slides is commonly known as a *counting semaphore*.
 - Since its value can range over an unrestricted domain.

Binary semaphore

- The semaphore construct described in the previous slides is commonly known as a *counting semaphore*.
 - Since its value can range over an unrestricted domain.
- A *binary semaphore* is a semaphore with an integer value range only between 0 and 1.

Binary semaphore

- The semaphore construct described in the previous slides is commonly known as a *counting semaphore*.
 - Since its value can range over an unrestricted domain.
- A *binary semaphore* is a semaphore with an integer value range only between 0 and 1.
 - It can be simpler to implement than a counting semaphore on some hardware architectures.

Binary semaphore

- The semaphore construct described in the previous slides is commonly known as a *counting semaphore*.
 - Since its value can range over an unrestricted domain.
- A *binary semaphore* is a semaphore with an integer value range only between 0 and 1.
 - It can be simpler to implement than a counting semaphore on some hardware architectures.
- A counting semaphore can be implemented using binary semaphores.

Implementation of binary semaphore

Implementation of binary semaphore

```
typedef struct {  
    bool flag;  
  
    /*processes blocked by this binary-semaphore*/  
    struct PCB *L;  
} binary-semaphore;
```

Implementation of binary semaphore

```
typedef struct {
    bool flag;

    /*processes blocked by this binary-semaphore*/
    struct PCB *L;
} binary-semaphore;

void bP(binary-semaphore *bS) {
    if(bS->flag == true)
        bS->flag = false;
    else {
        Add (curproc) to bS->L;

        (curproc)->state=WAITING;
        scheduler();
    }
}
```

Implementation of binary semaphore

```
typedef struct {  
    bool flag;  
  
    /*processes blocked by this binary-semaphore*/  
    struct PCB *L;  
} binary-semaphore;
```

```
void bP(binary-semaphore *bS) {  
    if(bS->flag == true)  
        bS->flag = false;  
    else {  
        Add (curproc) to bS->L;  
  
        (curproc)->state=WAITING;  
        scheduler();  
    }  
}
```

```
void bV(binary-semaphore *bS) {  
    if(bS->L is empty)  
        bS->flag = true;  
    else {  
        Remove a proc (A) from bS->L;  
  
        (A)->state=READY; //Wakeup (A)  
    }  
}
```

Implementation of binary semaphore

```
typedef struct {  
    bool flag;  
  
    /*processes blocked by this binary-semaphore*/  
    struct PCB *L;  
} binary-semaphore;
```

```
void bP(binary-semaphore *bS) {  
    if(bS->flag == true)  
        bS->flag = false;  
    else {  
        Add (curproc) to bS->L;  
  
        (curproc)->state=WAITING;  
        scheduler();  
    }  
}
```

```
void bV(binary-semaphore *bS) {  
    if(bS->L is empty)  
        bS->flag = true;  
    else {  
        Remove a proc (A) from bS->L;  
  
        (A)->state=READY; //Wakeup (A)  
    }  
}
```

- The *bP* and *bV* must be executed **atomically**.

Implement counting semaphore using binary semaphore

Implement counting semaphore using binary semaphore

```
typedef struct {  
    int value;  
  
    binary_semaphore bS1 = true, /*protect 'value'*/  
                    bS2 = false; /*synchronization*/  
} semaphore;
```

Implement counting semaphore using binary semaphore

```
typedef struct {  
    int value;  
  
    binary_semaphore bS1 = true, /*protect 'value'*/  
                    bS2 = false; /*synchronization*/  
} semaphore;
```

```
void P(semaphore *S)  
{  
    bP(&S->bS1);  
    S->value--;  
    if (S->value < 0) {  
        bV(&S->bS1);  
        bP(&S->bS2);  
    } else  
        bV(&S->bS1);  
}
```

Implement counting semaphore using binary semaphore

```
typedef struct {  
    int value;  
  
    binary_semaphore bS1 = true, /*protect 'value'*/  
                    bS2 = false; /*synchronization*/  
} semaphore;
```

```
void P(semaphore *S)  
{  
    bP(&S->bS1);  
    S->value--;  
    if (S->value < 0) {  
        bV(&S->bS1);  
        bP(&S->bS2);  
    } else  
        bV(&S->bS1);  
}
```

```
void V(semaphore *S)  
{  
    bP(&S->bS1);  
    S->value++;  
    if (S->value <= 0)  
        bV(&S->bS2);  
    bV(&S->bS1);  
}
```


Implement counting semaphore using binary semaphore

```
typedef struct {  
    int value;  
  
    binary_semaphore bS1 = true, /*protect 'value'*/  
                    bS2 = false; /*synchronization*/  
} semaphore;
```

```
void P(semaphore *S)  
{  
    bP(&S->bS1);  
    S->value--;  
    if (S->value < 0) {  
        bV(&S->bS1);  
        bP(&S->bS2);  
    } else  
        bV(&S->bS1);  
}
```

```
void V(semaphore *S)  
{  
    bP(&S->bS1);  
    S->value++;  
    if (S->value <= 0)  
        bV(&S->bS2);  
    bV(&S->bS1);  
}
```

- This implementation of P and V may **not** be executed **atomically**.

Questions

- Any questions?



Deadlock and starvation (1/2)

Deadlock and starvation (1/2)

- Although semaphores provide a convenient and effective mechanism for process synchronization, their incorrect use can still result in hard-to-detect errors.

Deadlock and starvation (1/2)

- Although semaphores provide a convenient and effective mechanism for process synchronization, their incorrect use can still result in hard-to-detect errors.
 - For example, suppose that two P s in the producer loop were reversed in order. That is,

Deadlock and starvation (1/2)

- Although semaphores provide a convenient and effective mechanism for process synchronization, their incorrect use can still result in hard-to-detect errors.
 - For example, suppose that two P s in the producer loop were reversed in order. That is,

```
do { /*producer*/  
    ...  
    produce an item;  
    ...  
    P(&mutex); /*XXX*/  
    P(&empty); /*XXX*/  
    ...  
    add the item to buffer;  
    ...  
    V(&mutex);  
    V(&full);  
} while(1);
```

Deadlock and starvation (1/2)

- Although semaphores provide a convenient and effective mechanism for process synchronization, their incorrect use can still result in hard-to-detect errors.
 - For example, suppose that two *Ps* in the producer loop were reversed in order. That is,

```
do { /*producer*/  
    ...  
    produce an item;  
    ...  
    P(&mutex); /*XXX*/  
    P(&empty); /*XXX*/  
    ...  
    add the item to buffer;  
    ...  
    V(&mutex);  
    V(&full);  
} while(1);
```

```
do { /*consumer*/  
    P(&full);  
    P(&mutex);  
    ...  
    remove an item from buffer;  
    ...  
    V(&mutex);  
    V(&empty);  
    ...  
    consume the item;  
    ...  
} while(1);
```

Deadlock and starvation (2/2)

Deadlock and starvation (2/2)

- So, what's a deadlock?

Deadlock and starvation (2/2)

- So, what's a deadlock?
 - *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*

Deadlock and starvation (2/2)

- So, what's a deadlock?
 - *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*
 - We will cover the deadlock in the next chapter.

Deadlock and starvation (2/2)

- So, what's a deadlock?
 - *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*
 - We will cover the deadlock in the next chapter.
- Another problem related to deadlocks is *starvation* or *indefinite blocking*.

Deadlock and starvation (2/2)

- So, what's a deadlock?
 - *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*
 - We will cover the deadlock in the next chapter.
- Another problem related to deadlocks is *starvation* or *indefinite blocking*.
 - A situation where processes wait indefinitely within the semaphore.

Deadlock and starvation (2/2)

- So, what's a deadlock?
 - *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*
 - We will cover the deadlock in the next chapter.
- Another problem related to deadlocks is *starvation* or *indefinite blocking*.
 - A situation where processes wait indefinitely within the semaphore.
 - For example, the previous solution to the 'readers-writers' problem may result in starvation.

Deadlock and starvation (2/2)

- So, what's a deadlock?
 - *A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.*
 - We will cover the deadlock in the next chapter.
- Another problem related to deadlocks is *starvation* or *indefinite blocking*.
 - A situation where processes wait indefinitely within the semaphore.
 - For example, the previous solution to the 'readers-writers' problem may result in starvation.
 - The starvation can be avoided by using a FCFS resource allocation policy.

Questions

- Any questions?



Outline

- 1 Race condition
- 2 Possible solutions to race condition
 - Software solutions
 - Hardware solutions
 - Conclusion
- 3 Semaphore
 - What's the semaphore?
 - Programming interfaces
 - Classic problems of synchronization
 - Binary semaphore
 - Deadlock and starvation
- 4 Monitor
 - Condition variables
 - Monitor solution to the dining-philosopher problem
 - Language support
- 5 Relationships of semaphore and monitor

High-level language synchronization constructs

High-level language synchronization constructs

- As you can see, one subtle error when using semaphores may result in race conditions, deadlocks or other unpredictable and irreproducible behavior.

High-level language synchronization constructs

- As you can see, one subtle error when using semaphores may result in race conditions, deadlocks or other unpredictable and irreproducible behavior.
- To make the life of programmers easier, several high-level language synchronization constructs have been introduced.

High-level language synchronization constructs

- As you can see, one subtle error when using semaphores may result in race conditions, deadlocks or other unpredictable and irreproducible behavior.
- To make the life of programmers easier, several high-level language synchronization constructs have been introduced.
 - *Monitors.*

High-level language synchronization constructs

- As you can see, one subtle error when using semaphores may result in race conditions, deadlocks or other unpredictable and irreproducible behavior.
- To make the life of programmers easier, several high-level language synchronization constructs have been introduced.
 - *Monitors.*
 - It's suggested by Brinch-Hansen in 1973.

High-level language synchronization constructs

- As you can see, one subtle error when using semaphores may result in race conditions, deadlocks or other unpredictable and irreproducible behavior.
- To make the life of programmers easier, several high-level language synchronization constructs have been introduced.
 - *Monitors.*
 - It's suggested by Brinch-Hansen in 1973.
 - And many others...

Monitor

Monitor

- What's a monitor?

- What's a monitor?
 - A monitor is a collection of procedures, variables and data structures that are all grouped together in a special kind of module or package.

- What's a monitor?
 - A monitor is a collection of procedures, variables and data structures that are all grouped together in a special kind of module or package.
- Monitors have an important property that makes them useful for achieving mutual exclusion: *only one process can be active in a monitor at any instant.*

Monitor

- What's a monitor?

- A monitor is a collection of procedures, variables and data structures that are all grouped together in a special kind of module or package.

```
monitor producer-consumer {  
    int i;  
    condition c;  
  
    void produce()  
    {  
        ...  
    }  
  
    void consume()  
    {  
        ...  
    }  
}
```

- Monitors have an important property that makes them useful for achieving mutual exclusion: *only one process can be active in a monitor at any instant.*

Condition variables

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.
- *Condition variables*

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.
- *Condition variables*
 - Condition variable is defined as a type of variables, which is associated with **only** two operations:

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.
- *Condition variables*
 - Condition variable is defined as a type of variables, which is associated with **only** two operations:
 - *c.wait()*;

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.
- *Condition variables*
 - Condition variable is defined as a type of variables, which is associated with **only** two operations:
 - *c.wait()*; - means that the process invoking it will be suspended until another process invokes

Condition variables


- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.
- *Condition variables*
 - Condition variable is defined as a type of variables, which is associated with **only** two operations:
 - *c.wait()*; - means that the process invoking it will be suspended until another process invokes
 - *c.signal()*;

Condition variables

- The monitor introduced so far is not enough for a reasonable solution.
 - We also need a way for processes to block when they cannot proceed within a monitor.
- *Condition variables*
 - Condition variable is defined as a type of variables, which is associated with **only** two operations:
 - *c.wait()*; - means that the process invoking it will be suspended until another process invokes
 - *c.signal()*;
 - Note that if no process is suspended, then the *signal* operation has no effect.

Monitor with condition variables

Monitor with condition variables



What happens after a `c.signal()`?

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.
- After *P* completed the *c.signal()*, it's possible that both *P* and *Q* will be active simultaneously within the monitor.

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.
- After *P* completed the *c.signal()*, it's possible that both *P* and *Q* will be active simultaneously within the monitor.
 - This will break the property of the monitor!

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.
- After *P* completed the *c.signal()*, it's possible that both *P* and *Q* will be active simultaneously within the monitor.
 - This will break the property of the monitor!
- Three possibilities exist:

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.
- After *P* completed the *c.signal()*, it's possible that both *P* and *Q* will be active simultaneously within the monitor.
 - This will break the property of the monitor!
- Three possibilities exist:
 - Hoare-style: suspending *P* and letting *Q* run.

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.
- After *P* completed the *c.signal()*, it's possible that both *P* and *Q* will be active simultaneously within the monitor.
 - This will break the property of the monitor!
- Three possibilities exist:
 - Hoare-style: suspending *P* and letting *Q* run.
 - Brinch-Hansen-style: *P* must leave the monitor immediately.

What happens after a *c.signal()*?

- Suppose a process *P* is invoking *c.signal()* and another process *Q* is blocked by the condition variable *c*.
- After *P* completed the *c.signal()*, it's possible that both *P* and *Q* will be active simultaneously within the monitor.
 - This will break the property of the monitor!
- Three possibilities exist:
 - Hoare-style: suspending *P* and letting *Q* run.
 - Brinch-Hansen-style: *P* must leave the monitor immediately.
 - Mesa-style (Mesa is a programming language): letting *P* run and suspending *Q*.

Monitor solution to the dining-philosopher problem

Monitor solution to the dining-philosopher problem

```
monitor dp {
    enum {THINKING, HUNGRY, EATING} state[N];
    condition c[N];

    void take_forks(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            c[i].wait();
    }

    void put_forks(int i) {
        state[i] = THINKING;
        test(LEFT);
        test(RIGHT);
    }

    void test(int i) {
        if (state[i] == HUNGRY &&
            state[LEFT] != EATING &&
            state[RIGHT] != EATING) {
            state[i] = EATING;
            c[i].signal();
        }
    }

    void init() {
        for (int i = 0; i < N; i++)
            state[i] = THINKING;
    }
}
```

Questions

- Any questions?



Language support

Language support

- The monitor constructs must be supported by the programming language to be useful.

Language support

- The monitor constructs must be supported by the programming language to be useful.
 - That is, the compiler must recognize the monitor construct and generate codes to support its functionality.

Language support

- The monitor constructs must be supported by the programming language to be useful.
 - That is, the compiler must recognize the monitor construct and generate codes to support its functionality.
- Example: Java (Mesa-style monitor with only one condition variable)

Language support

- The monitor constructs must be supported by the programming language to be useful.
 - That is, the compiler must recognize the monitor construct and generate codes to support its functionality.
- Example: Java (Mesa-style monitor with only one condition variable)
 - By adding the keyword *synchronized* to a method declaration, Java guarantees that once any thread has started executing that method, no other thread will be allowed to start executing any other *synchronized* method in that class.

- The monitor constructs must be supported by the programming language to be useful.
 - That is, the compiler must recognize the monitor construct and generate codes to support its functionality.
- Example: Java (Mesa-style monitor with only one condition variable)
 - By adding the keyword *synchronized* to a method declaration, Java guarantees that once any thread has started executing that method, no other thread will be allowed to start executing any other *synchronized* method in that class.
 - And Java provides two operations: *wait* and *notify* to block and wakeup the thread.

Questions

- Any questions?



Outline

- 1 Race condition
- 2 Possible solutions to race condition
 - Software solutions
 - Hardware solutions
 - Conclusion
- 3 Semaphore
 - What's the semaphore?
 - Programming interfaces
 - Classic problems of synchronization
 - Binary semaphore
 - Deadlock and starvation
- 4 Monitor
 - Condition variables
 - Monitor solution to the dining-philosopher problem
 - Language support
- 5 Relationships of semaphore and monitor

Relationship of semaphore and monitor

Relationship of semaphore and monitor

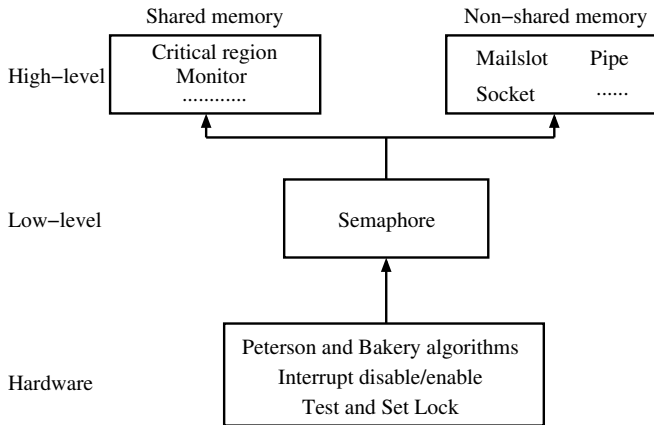
- Semaphore and monitor are *equivalent* in their functionality.

Relationship of semaphore and monitor

- Semaphore and monitor are *equivalent* in their functionality.
 - But their use and implementation are quite different.

Relationship of semaphore and monitor

- Semaphore and monitor are *equivalent* in their functionality.
 - But their use and implementation are quite different.



Implement semaphore using monitor

Implement semaphore using monitor

```
monitor semaphore {  
    int value;  
    condition c;  
  
    void P() {  
        value--;  
        if (value < 0)  
            c.wait();  
    }  
  
    void V() {  
        value++;  
        if (value <= 0)  
            c.signal();  
    }  
}
```


Implement monitor using semaphore

Implement monitor using semaphore

```
semaphore mutex = 1, next = 0;  
int next_count = 0;  
  
// for each condition variable x  
semaphore x_sem = 1;  
int x_count = 0;
```

Implement monitor using semaphore

```
semaphore mutex = 1, next = 0;
int next_count = 0;

// for each condition variable x
semaphore x_sem = 1;
int x_count = 0;

// Mutual exclusion
// within a monitor
P(&mutex);
...
body of F;
...
if(next_count > 0)
    V(&next);
else
    V(&mutex);
```

Implement monitor using semaphore

```
semaphore mutex = 1, next = 0;
int next_count = 0;

// for each condition variable x
semaphore x_sem = 1;
int x_count = 0;

// Mutual exclusion
// within a monitor
P(&mutex);
...
body of F;
...
if(next_count > 0)
    V(&next);
else
    V(&mutex);

// x.wait()
x_count++;
if(next_count > 0)
    V(&next);
else
    V(&mutex);
P(&x_sem);
x_count--;
```

Implement monitor using semaphore

```
semaphore mutex = 1, next = 0;
int next_count = 0;

// for each condition variable x
semaphore x_sem = 1;
int x_count = 0;

// Mutual exclusion
// within a monitor
P(&mutex);
...
body of F;
...
if (next_count > 0)
    V(&next);
else
    V(&mutex);

// x.wait()
x_count++;
if (next_count > 0)
    V(&next);
else
    V(&mutex);
P(&x_sem);
x_count--;

// x.signal()
if (x_count > 0) {
    next_count++;
    V(&x_sem);
    P(&next);
    next_count--;
}
```

Questions

- Any questions?

