

操作系统原理

第五章：线程管理

洪明坚

重庆大学软件学院

February 19, 2016

1 Thread

- Process revisited
- What's thread?
- Single- and multi-threaded process
- Benefits of thread

2 Thread implementation

- User thread
- Kernel thread

3 Multi-threaded programming

- Multi-threaded APIs
- Pthreads
- Win32 threads

Outline

1 Thread

- Process revisited
- What's thread?
- Single- and multi-threaded process
- Benefits of thread

2 Thread implementation

- User thread
- Kernel thread

3 Multi-threaded programming

- Multi-threaded APIs
- Pthreads
- Win32 threads

Process revisited

Process revisited

- The process is

Process revisited

- The process is
 - a unit of resource allocation;

Process revisited

- The process is
 - a unit of resource allocation;
 - a unit of dispatching (scheduling).

Process revisited

- The process is
 - a unit of resource allocation;
 - a unit of dispatching (scheduling).
- Traditionally, a process has only one thread of control.

Process revisited

- The process is
 - a unit of resource allocation;
 - a unit of dispatching (scheduling).
- Traditionally, a process has only one thread of control.
 - If we separate the above two concepts and allow multiple threads of control within one process, we get the threads.

Process revisited

- The process is
 - a unit of resource allocation;
 - a unit of dispatching (scheduling).
- Traditionally, a process has only one thread of control.
 - If we separate the above two concepts and allow multiple threads of control within one process, we get the threads.
 - That is, processes are used to group resources together;

Process revisited

- The process is
 - a unit of resource allocation;
 - a unit of dispatching (scheduling).
- Traditionally, a process has only one thread of control.
 - If we separate the above two concepts and allow multiple threads of control within one process, we get the threads.
 - That is, processes are used to group resources together; threads are the entities dispatched (scheduled) for execution on the CPU.

Thread (1/2)

Thread (1/2)

- A thread is a basic unit of CPU utilization in modern operating system.

Thread (1/2)

- A thread is a basic unit of CPU utilization in modern operating system.
 - Also known as **lightweight process (LWP)**.

Thread (1/2)

- A thread is a basic unit of CPU utilization in modern operating system.
 - Also known as **lightweight process (LWP)**.
- **Multi-threading**

Thread (1/2)

- A thread is a basic unit of CPU utilization in modern operating system.
 - Also known as **lightweight process (LWP)**.
- **Multi-threading**
 - Allowing multiple threads in the same process.

Thread (1/2)

- A thread is a basic unit of CPU utilization in modern operating system.
 - Also known as **lightweight process (LWP)**.
- **Multi-threading**
 - Allowing multiple threads in the same process.
 - They share resources belonging to the same process, such as its code section, data section, open files, etc.

Thread (1/2)

- A thread is a basic unit of CPU utilization in modern operating system.
 - Also known as **lightweight process (LWP)**.
- **Multi-threading**
 - Allowing multiple threads in the same process.
 - They share resources belonging to the same process, such as its code section, data section, open files, etc.
 - But, each thread within one process has a **private thread context (including the CPU register set and other state information) and a private stack**.

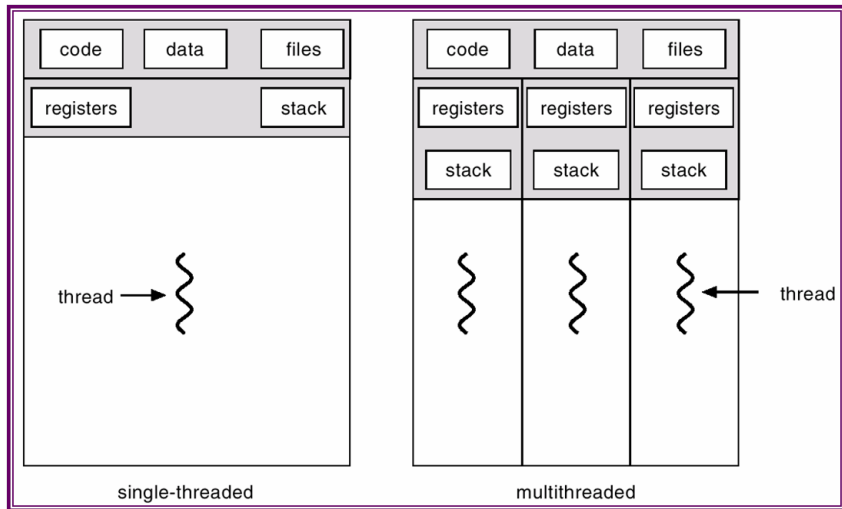
Thread (2/2)

Thread (2/2)

Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

Single- and multi-threaded process

Single- and multi-threaded process



Benefits of thread

Benefits of thread

- Responsiveness

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing
 - Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel.

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing
 - Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel.
- Economy

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing
 - Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel.
- Economy
 - Takes MUCH less time and resource to create a new thread than a process.

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing
 - Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel.
- Economy
 - Takes MUCH less time and resource to create a new thread than a process.
 - Takes MUCH less time to context switch threads within the same process.

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing
 - Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel.
- Economy
 - Takes MUCH less time and resource to create a new thread than a process.
 - Takes MUCH less time to context switch threads within the same process.
- Utilization of multiprocessor architectures

Benefits of thread

- Responsiveness
 - Allows other threads to continue responding to the user even if one or several threads is blocked or is performing a lengthy operation.
- Resource sharing
 - Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel.
- Economy
 - Takes MUCH less time and resource to create a new thread than a process.
 - Takes MUCH less time to context switch threads within the same process.
- Utilization of multiprocessor architectures
 - Parallelism is possible by assigning each CPU a thread.

Questions

- Any questions?



Outline

- 1 Thread
 - Process revisited
 - What's thread?
 - Single- and multi-threaded process
 - Benefits of thread
- 2 Thread implementation
 - User thread
 - Kernel thread
- 3 Multi-threaded programming
 - Multi-threaded APIs
 - Pthreads
 - Win32 threads

Thread implementation

Thread implementation

- Multi-threading can be implemented

Thread implementation

- Multi-threading can be implemented
 - at user space for **user threads**;

Thread implementation

- Multi-threading can be implemented
 - at user space for **user threads**;
 - at the kernel for **kernel threads**.

Thread implementation

- Multi-threading can be implemented
 - at user space for **user threads**;
 - at the kernel for **kernel threads**.
 - with a hybrid scheme by combining user- and kernel- threads.

User threads (1/2)

User threads (1/2)

- It's implemented outside of the kernel as a **thread library** at the user space.

User threads (1/2)

- It's implemented outside of the kernel as a **thread library** at the user space.
 - It's the library which provides support for thread creation, scheduling and management.

User threads (1/2)

- It's implemented outside of the kernel as a **thread library** at the user space.
 - It's the library which provides support for thread creation, scheduling and management.
 - As far as the kernel is concerned, it's managing ordinary, single-threaded process.

User threads (1/2)

- It's implemented outside of the kernel as a **thread library** at the user space.
 - It's the library which provides support for thread creation, scheduling and management.
 - As far as the kernel is concerned, it's managing ordinary, single-threaded process.
- Examples

User threads (1/2)

- It's implemented outside of the kernel as a **thread library** at the user space.
 - It's the library which provides support for thread creation, scheduling and management.
 - As far as the kernel is concerned, it's managing ordinary, single-threaded process.
- Examples
 - Mach **C-threads**

User threads (1/2)

- It's implemented outside of the kernel as a **thread library** at the user space.
 - It's the library which provides support for thread creation, scheduling and management.
 - As far as the kernel is concerned, it's managing ordinary, single-threaded process.
- Examples
 - Mach **C-threads**
 - Solaris 2 **UI-threads**

User threads (2/2)

User threads (2/2)

- Advantages of the user threads

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.
 - This will save a lot of CPU cycles.

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.
 - This will save a lot of CPU cycles.
 - Allow each process to have its own customized scheduling algorithm.

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.
 - This will save a lot of CPU cycles.
 - Allow each process to have its own customized scheduling algorithm.
- Disadvantages of the user threads

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.
 - This will save a lot of CPU cycles.
 - Allow each process to have its own customized scheduling algorithm.
- Disadvantages of the user threads
 - Any user-level thread performing a blocking system call will cause the entire process to block.

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.
 - This will save a lot of CPU cycles.
 - Allow each process to have its own customized scheduling algorithm.
- Disadvantages of the user threads
 - Any user-level thread performing a blocking system call will cause the entire process to block.
 - Even if other threads are ready to run within the process.

User threads (2/2)

- Advantages of the user threads
 - Thread management and context switch need not to trap to the kernel.
 - This will save a lot of CPU cycles.
 - Allow each process to have its own customized scheduling algorithm.
- Disadvantages of the user threads
 - Any user-level thread performing a blocking system call will cause the entire process to block.
 - Even if other threads are ready to run within the process.
 - On a system with multiprocessors, the user-level threads cannot be dispatched for execution in parallel.

Kernel threads

Kernel threads

- It's supported directly by the operating system.

Kernel threads

- It's supported directly by the operating system.
 - The kernel performs thread creation, scheduling and management in kernel space.

Kernel threads

- It's supported directly by the operating system.
 - The kernel performs thread creation, scheduling and management in kernel space.
- Examples

Kernel threads

- It's supported directly by the operating system.
 - The kernel performs thread creation, scheduling and management in kernel space.
- Examples
 - Windows NT/XP

Kernel threads

- It's supported directly by the operating system.
 - The kernel performs thread creation, scheduling and management in kernel space.
- Examples
 - Windows NT/XP
 - Solaris

Kernel threads

- It's supported directly by the operating system.
 - The kernel performs thread creation, scheduling and management in kernel space.
- Examples
 - Windows NT/XP
 - Solaris
- Advantages and disadvantages

Kernel threads

- It's supported directly by the operating system.
 - The kernel performs thread creation, scheduling and management in kernel space.
- Examples
 - Windows NT/XP
 - Solaris
- Advantages and disadvantages
 - Inverse of the user threads.

Questions

- Any questions?



Outline

1 Thread

- Process revisited
- What's thread?
- Single- and multi-threaded process
- Benefits of thread

2 Thread implementation

- User thread
- Kernel thread

3 Multi-threaded programming

- Multi-threaded APIs
- Pthreads
- Win32 threads

Multi-threaded application programming interface

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Multi-threaded application programming interface

- **Pthreads**

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Multi-threaded application programming interface

- **Pthreads**

- Pthreads refers to the POSIX¹ standard (IEEE 1003.1c) defining an API for thread creation and synchronization.

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Multi-threaded application programming interface

• Pthreads

- Pthreads refers to the POSIX¹ standard (IEEE 1003.1c) defining an API for thread creation and synchronization.
 - It's a *specification* for thread behavior, NOT an *implementation*.

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Multi-threaded application programming interface

- **Pthreads**

- Pthreads refers to the POSIX¹ standard (IEEE 1003.1c) defining an API for thread creation and synchronization.
 - It's a *specification* for thread behavior, NOT an *implementation*.

- Win32

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Multi-threaded application programming interface

- **Pthreads**

- Pthreads refers to the POSIX¹ standard (IEEE 1003.1c) defining an API for thread creation and synchronization.
 - It's a *specification* for thread behavior, NOT an *implementation*.

- **Win32**

- CreateThread, ExitThread and TerminateThread, etc.

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Multi-threaded application programming interface

- **Pthreads**

- Pthreads refers to the POSIX¹ standard (IEEE 1003.1c) defining an API for thread creation and synchronization.
 - It's a *specification* for thread behavior, NOT an *implementation*.

- **Win32**

- CreateThread, ExitThread and TerminateThread, etc.
- POSIX Threads for Win32 (<http://sources.redhat.com/pthreads-win32>)

¹POSIX or “Portable Operating System Interface for uniX” is the collective name of a family of related standards specified by the IEEE to define the application programming interface (API) for software compatible with variants of the Unix operating system.

Example: Pthreads

Example: Pthreads

```
#include <pthread.h>
```

```
int sum; /*shared variable*/
```

```
/*the thread function*/
```

```
void *runner(void *arg);
```

```
int main()
```

```
{
```

```
pthread_t tid; /*thread ID*/
```

```
pthread_attr_t attr; /*attributes*/
```

```
/*get default attributes*/
```

```
pthread_attr_init(&attr);
```

```
/*create the thread*/
```

```
pthread_create(&tid, &attr, runner, 82);
```

```
/*wait for the thread to exit*/
```

```
pthread_join(tid, NULL);
```

```
printf("sum = %d\n", sum);
```

```
}
```

```
void *runner(void *arg)
{
    int i, upper=(int) arg;
    sum=0;
    if(upper>0)
        for(i=1; i<=upper; i++)
            sum+=i;

    pthread_exit(0);
    //or return (void *)0;
}
```

Example: Win32

Example: Win32

```
#include <windows.h>

int sum; /*shared variable*/

/*the thread function*/
DWORD WINAPI runner(LPVOID arg);

int main()
{
    DWORD tid; /*the thread ID*/
    HANDLE hThr; /*the thread handle*/

    /*create the thread*/
    hThr=CreateThread(0,0,runner,82,0,&tid);

    /*wait for the thread to exit*/
    WaitForSingleObject(hThr,INFINITE);

    printf("sum=%d\n",sum);
}
```

```
DWORD WINAPI runner(LPVOID
    arg)
{
    int i, upper = (int)arg
    ;
    sum = 0;
    if(upper > 0)
        for(i = 1; i <=
            upper; i++)
            sum += i;

    ExitThread(0);
    // or return 0L;
}
```

Questions

- Any questions?

